

# Introduction to Natural Language Processing

---

Matthieu Labeau

[matthieu.labeau@telecom-paris.fr](mailto:matthieu.labeau@telecom-paris.fr)

# Language Processing: goals

---

Interdisciplinary field, whose goal is to get computers to perform useful tasks [...] like enabling human-machine communication, improving human-human communication, or simply doing useful processing of text or speech.

Speech and Language Processing (Jurafsky & Manning) (Chapter I)

# Language Processing: goals

---

Interdisciplinary field, whose goal is to get computers to perform useful tasks [...] like enabling human-machine communication, improving human-human communication, or simply doing useful processing of text or speech.

Speech and Language Processing (Jurafsky & Manning) (Chapter I)

## Applications ?

From the same source: example of HAL 9000, in *2001: A Space Odyssey*  
→ Conversational agent - what does HAL imply ?

# Language Processing: goals

---

Interdisciplinary field, whose goal is to get computers to perform useful tasks [...] like enabling human-machine communication, improving human-human communication, or simply doing useful processing of text or speech.

Speech and Language Processing (Jurafsky & Manning) (Chapter I)

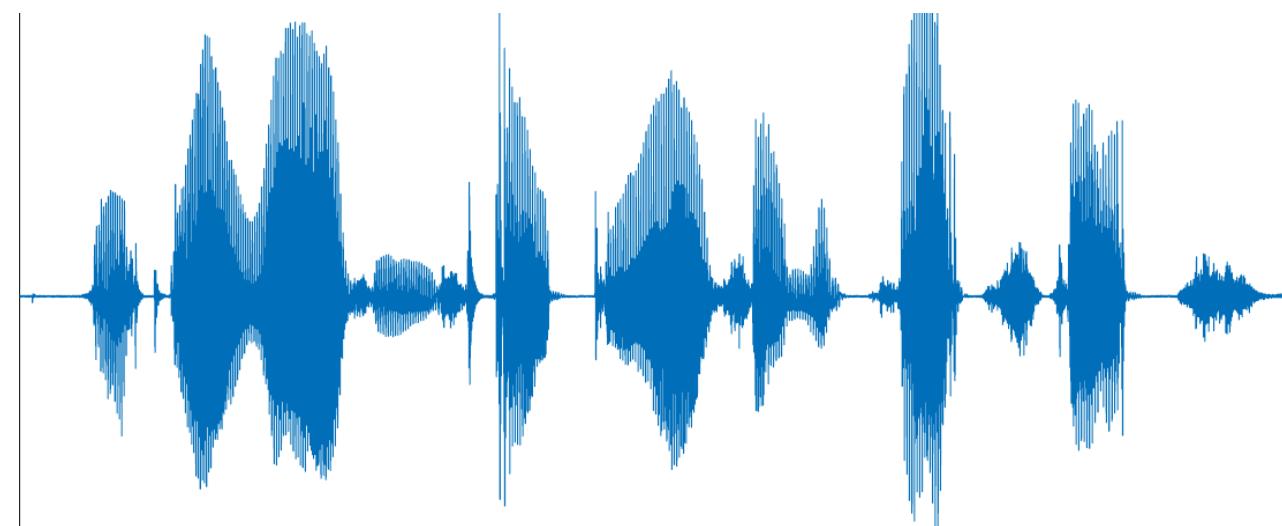
## Applications ?

From the same source: example of HAL 9000, in *2001: A Space Odyssey*  
→ Conversational agent - what does HAL imply ?

- Language input: *speech recognition, language understanding*
- Language output: *dialogue planning, speech synthesis*
- Information *retrieval, extraction*, and doing *inference* from it

# Analyzing language: tasks involved

---



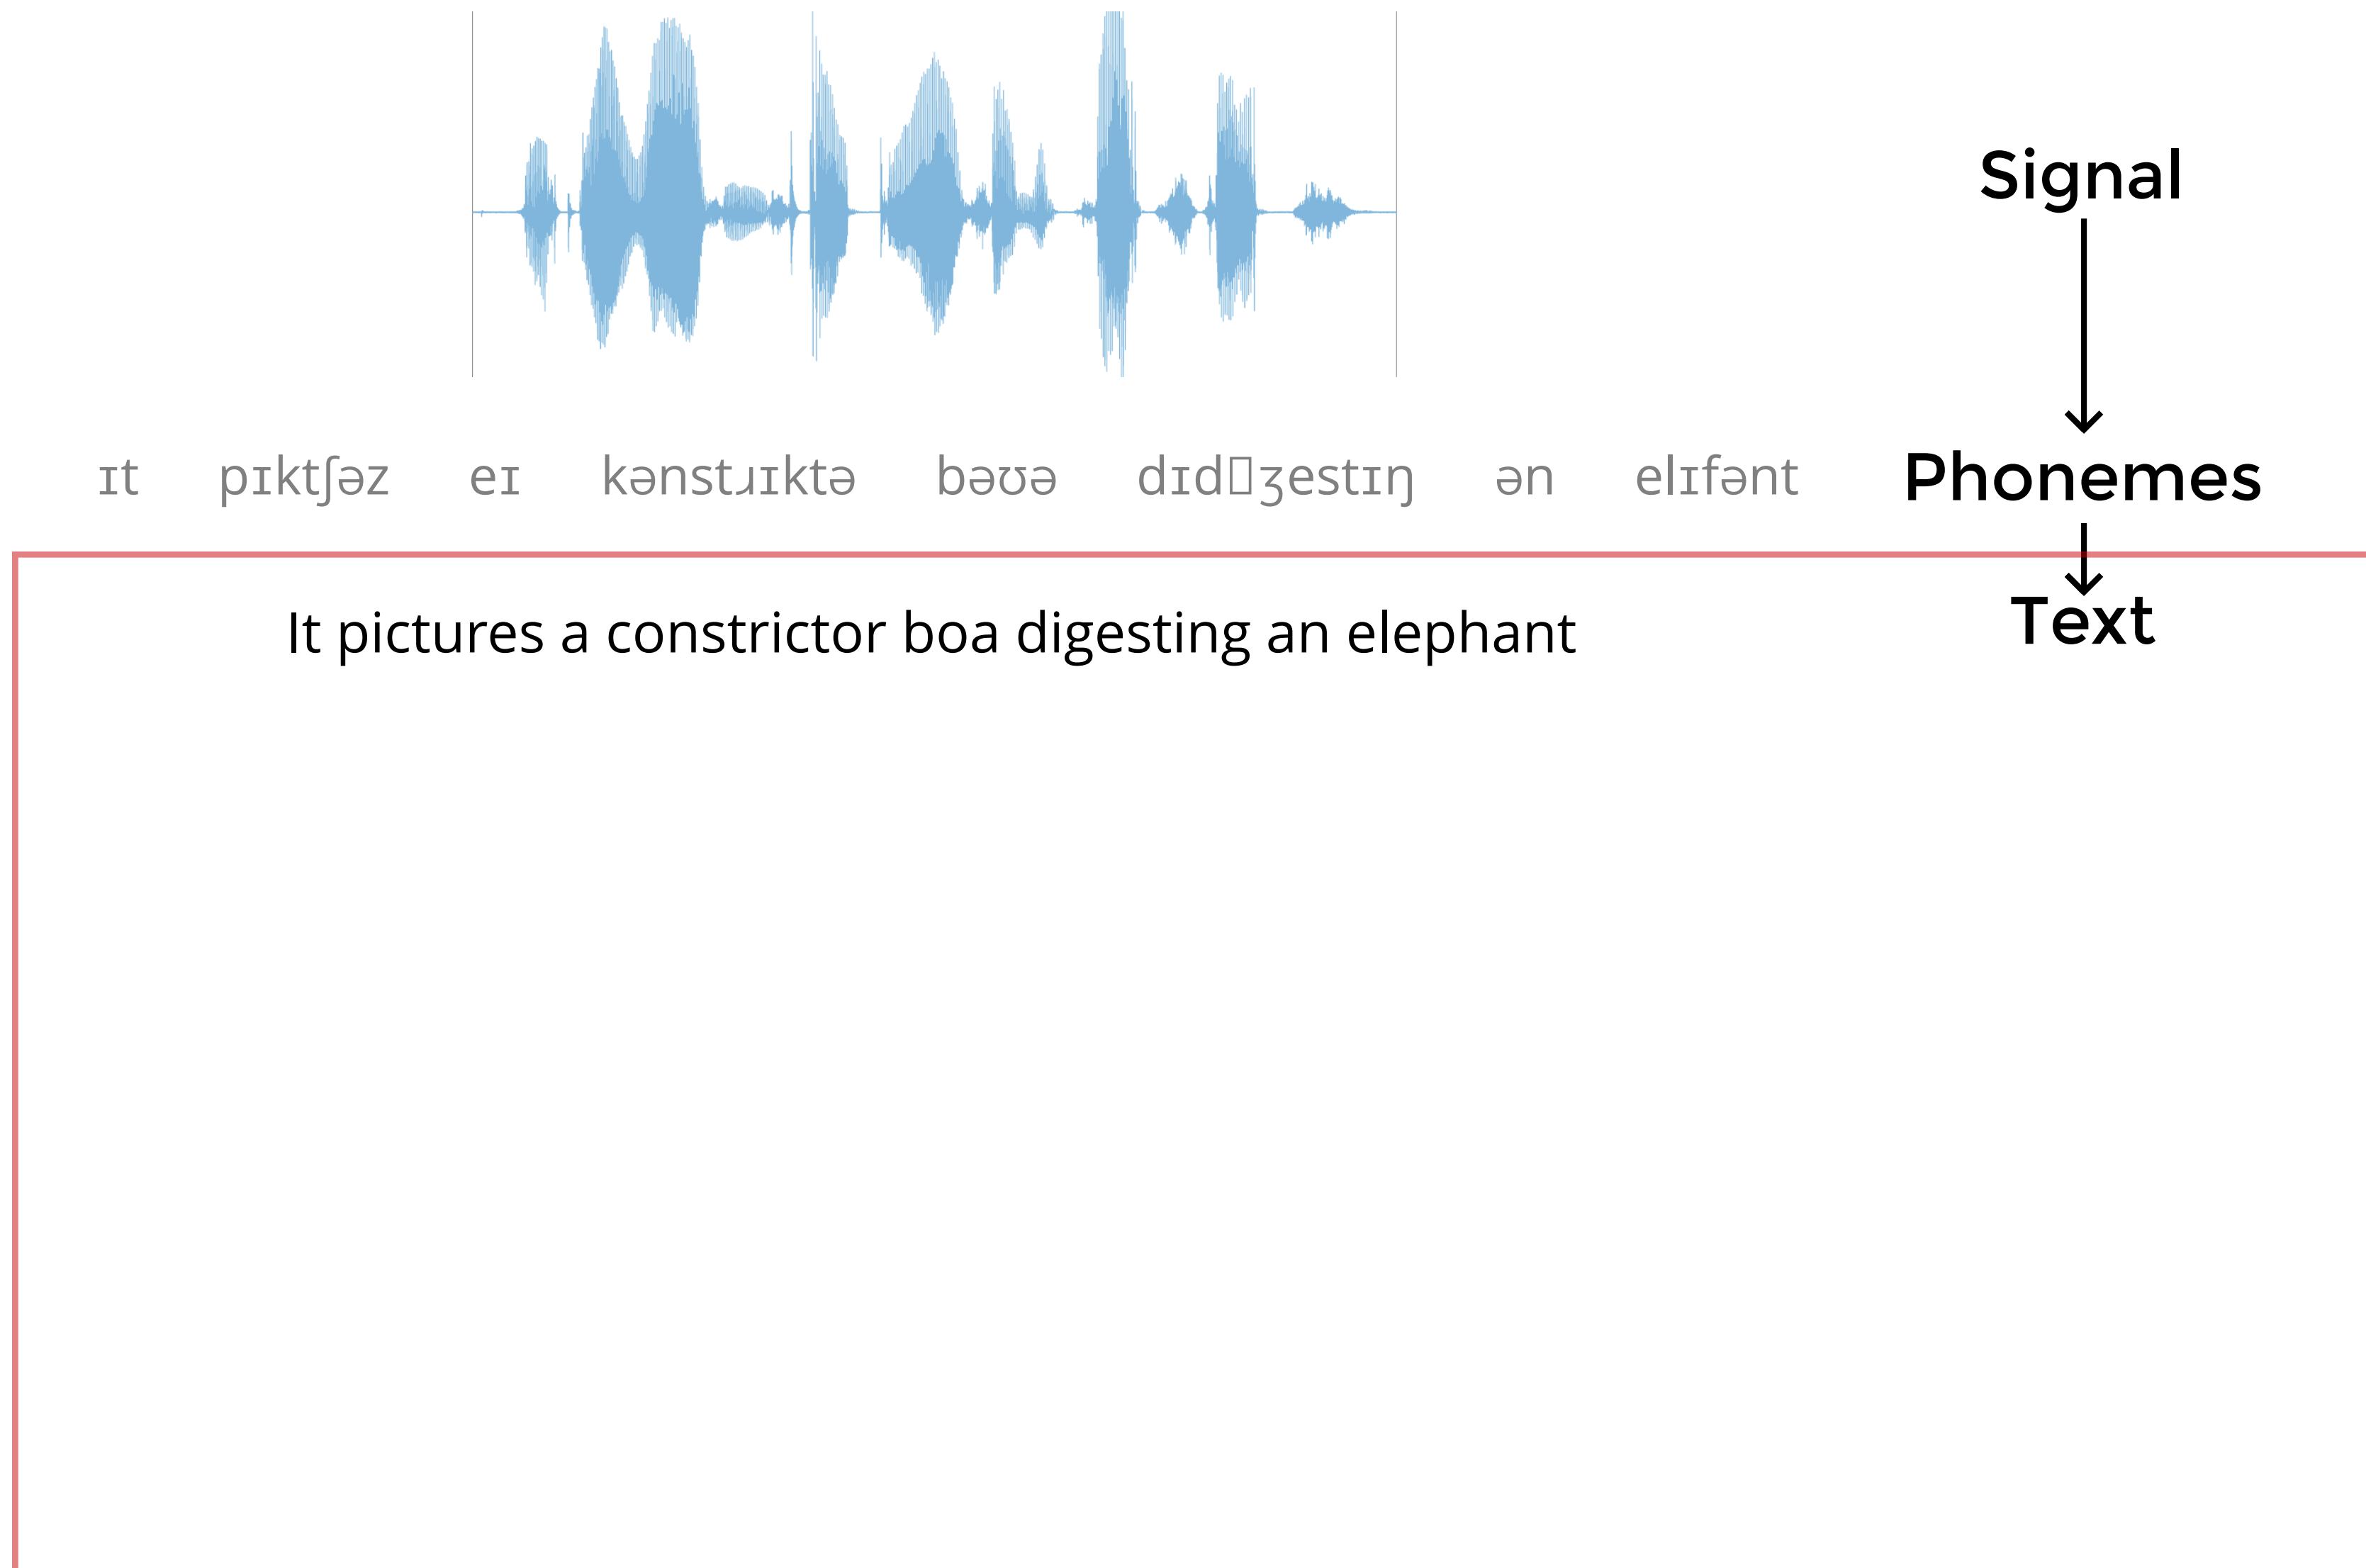
Signal



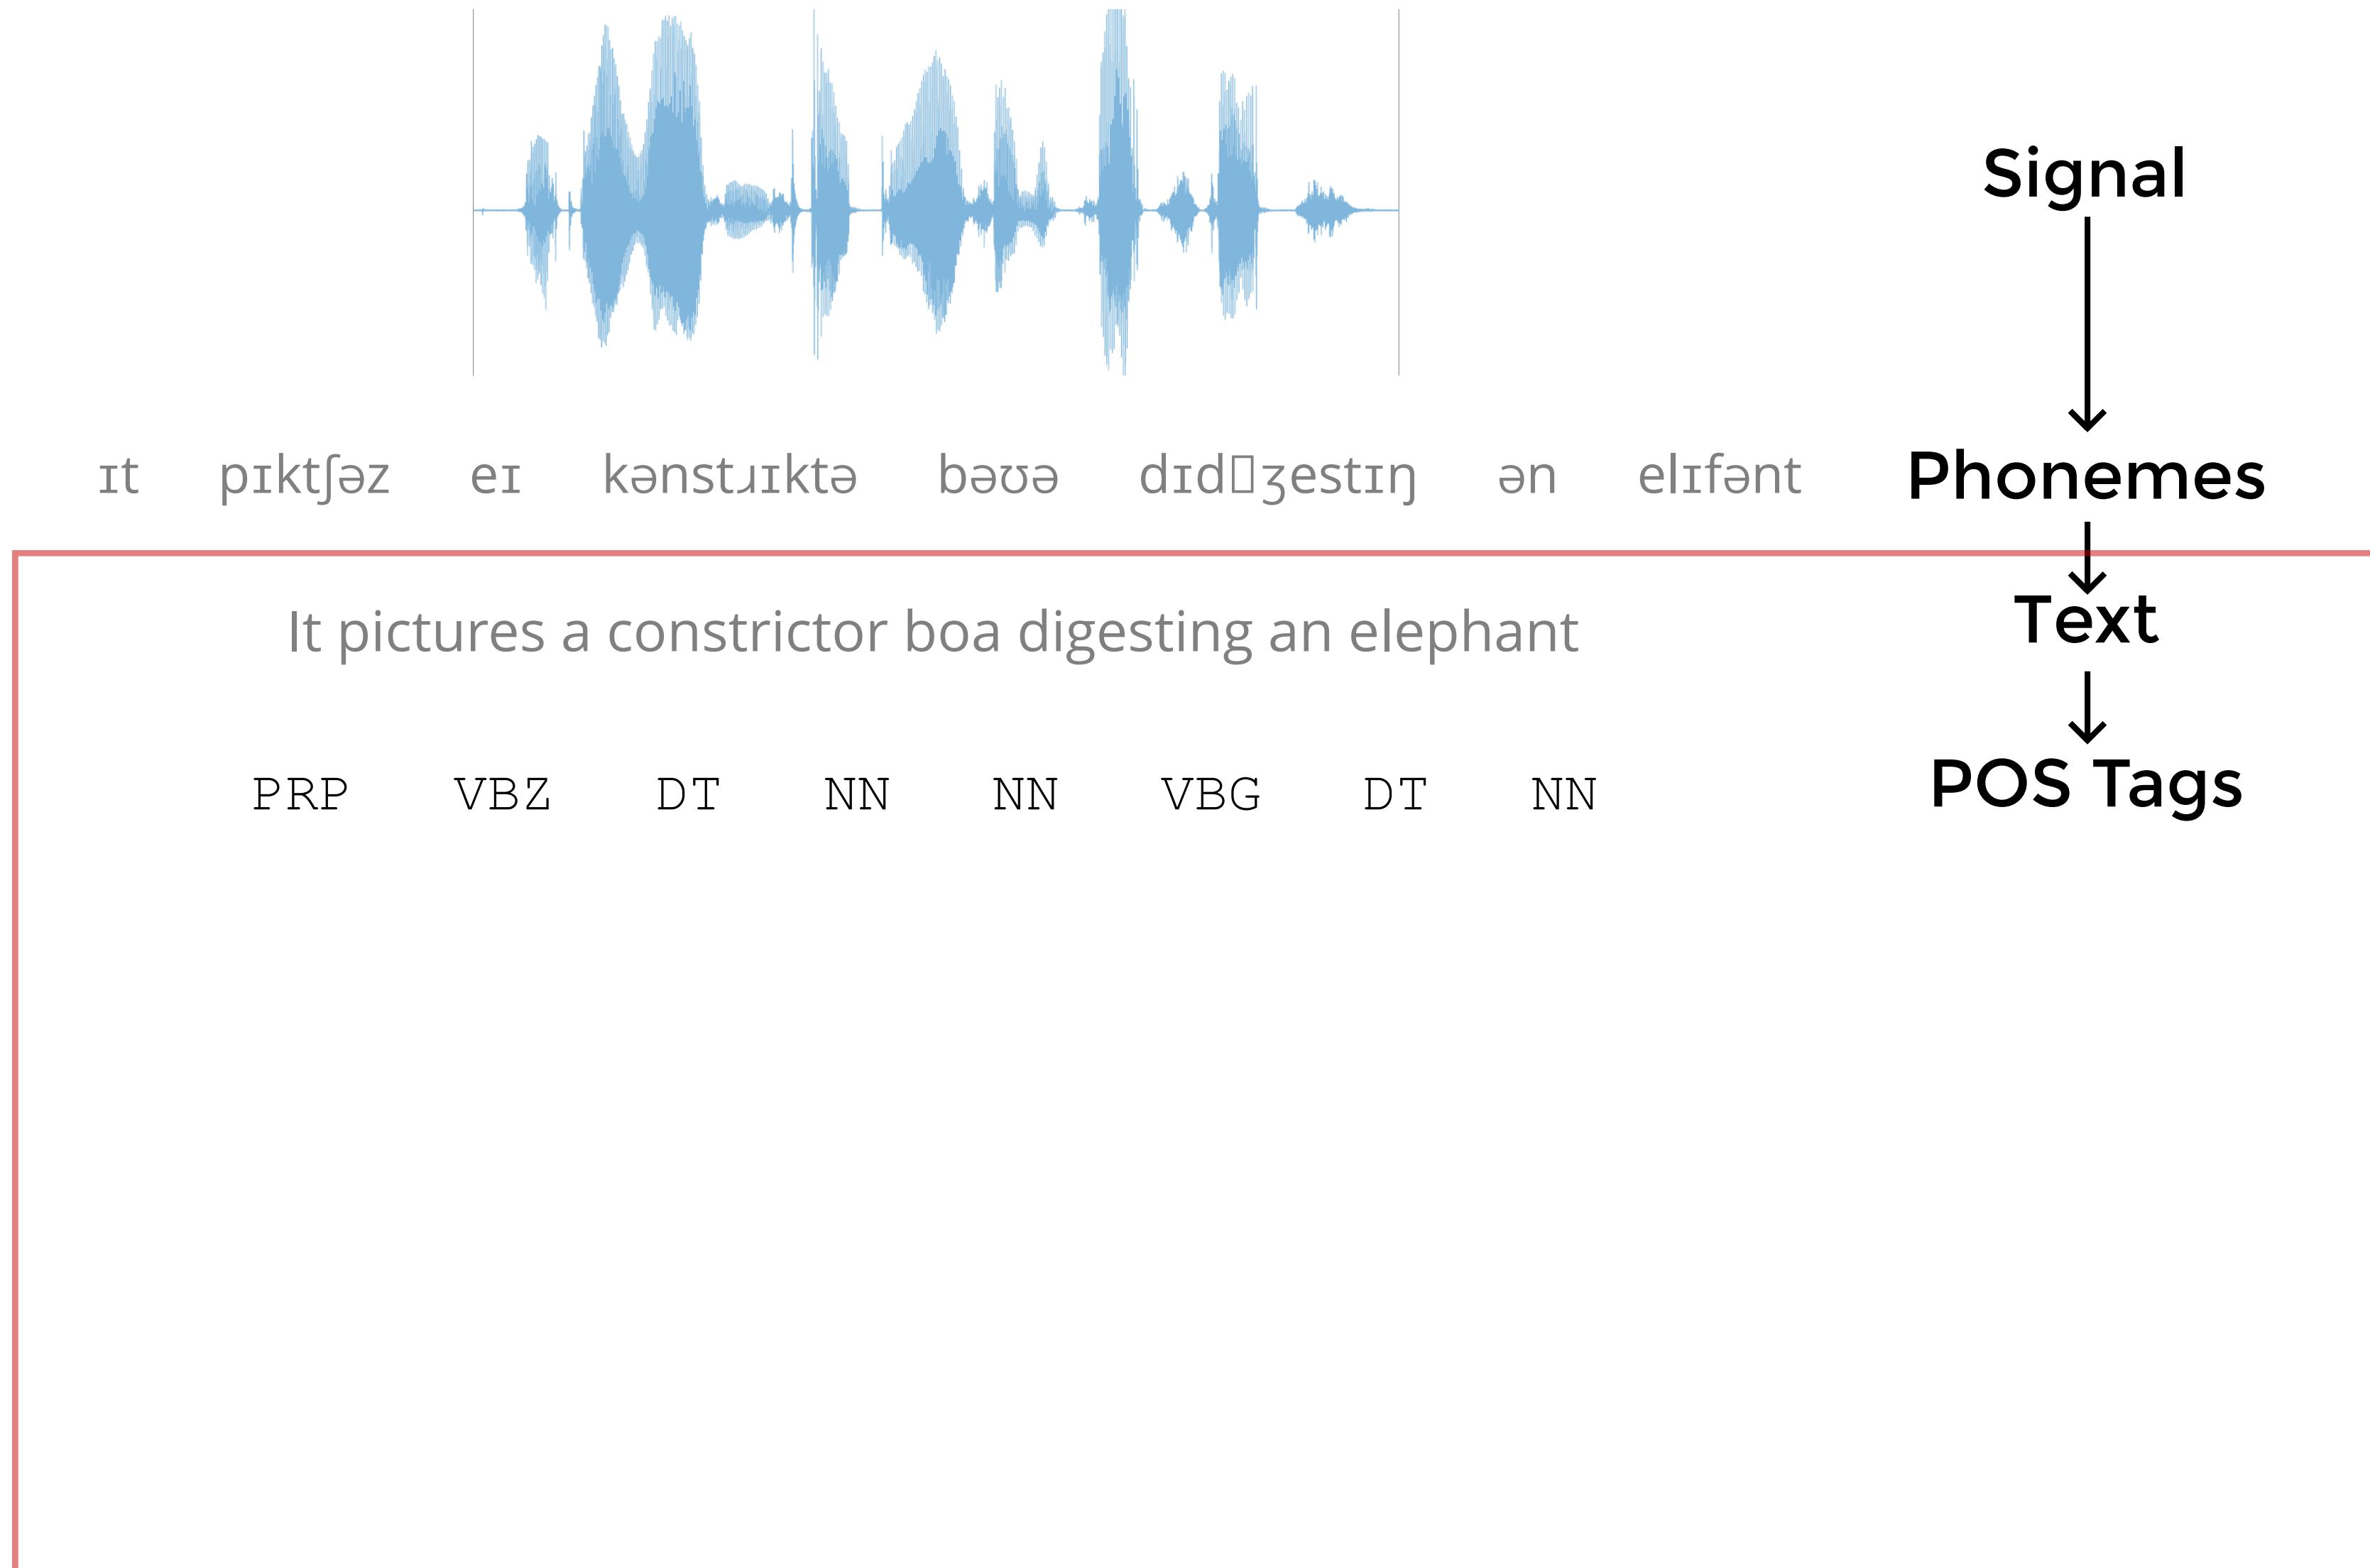
# Analyzing language: tasks involved



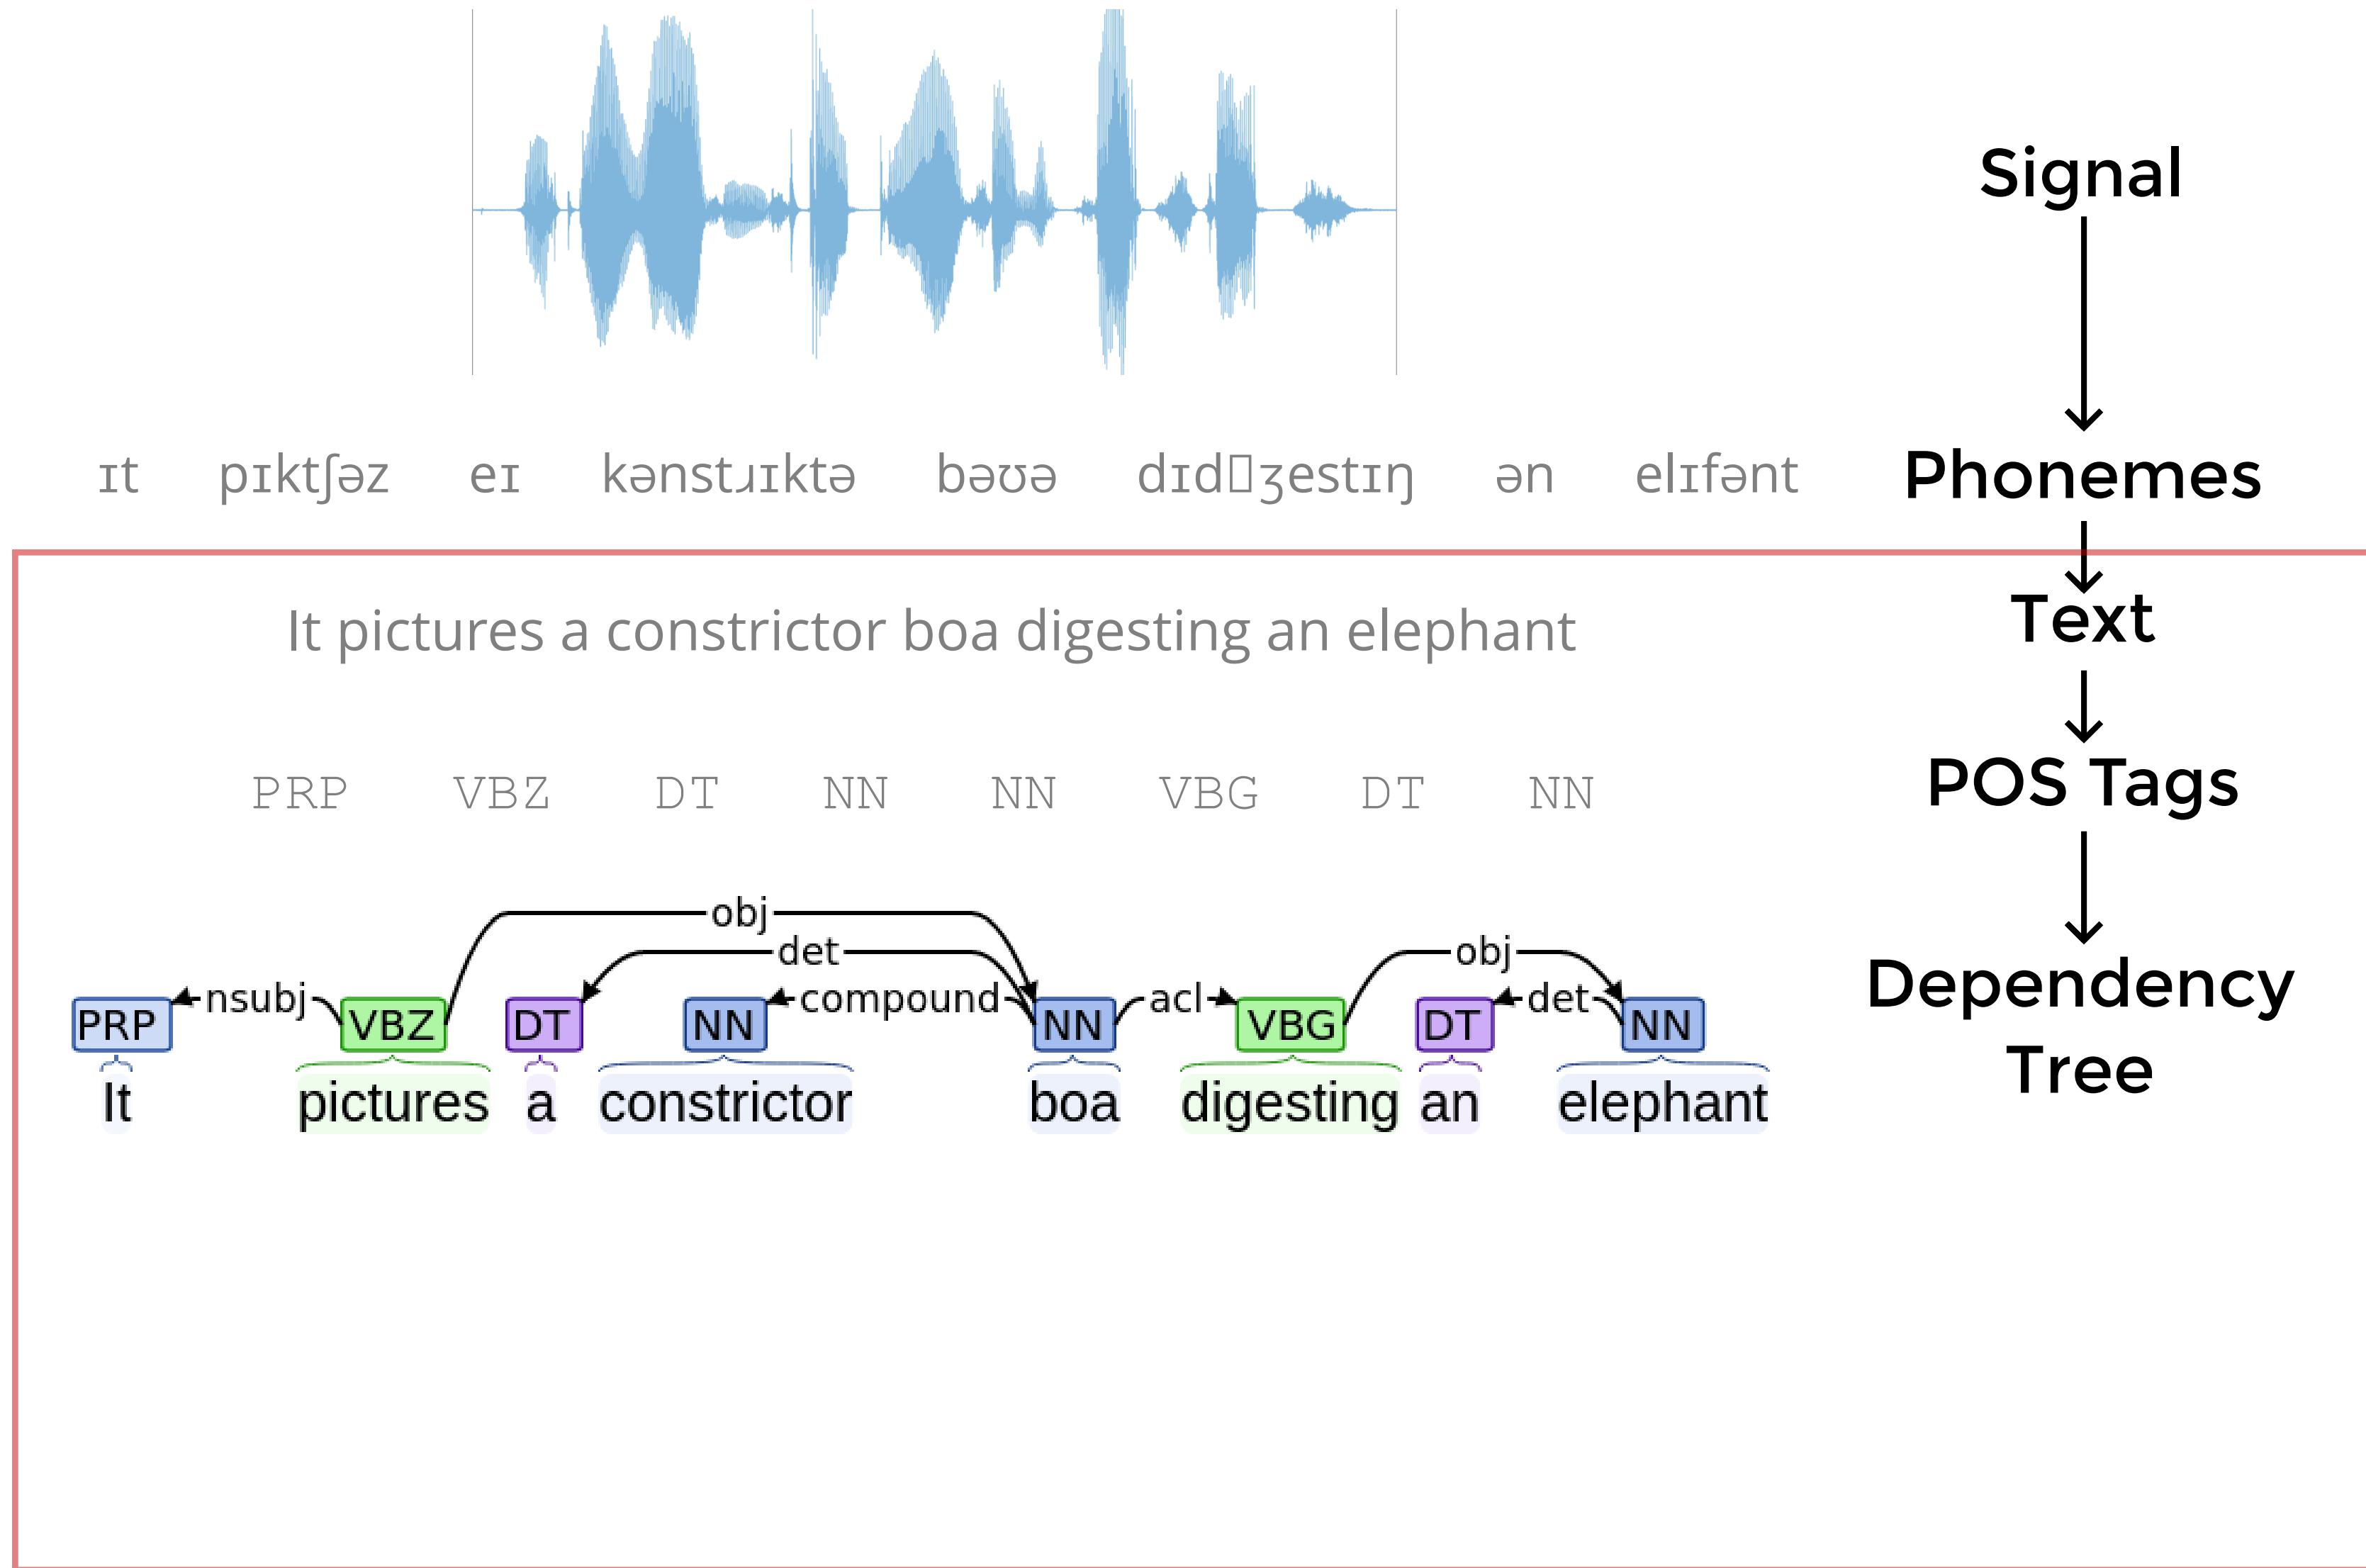
# Analyzing language: tasks involved



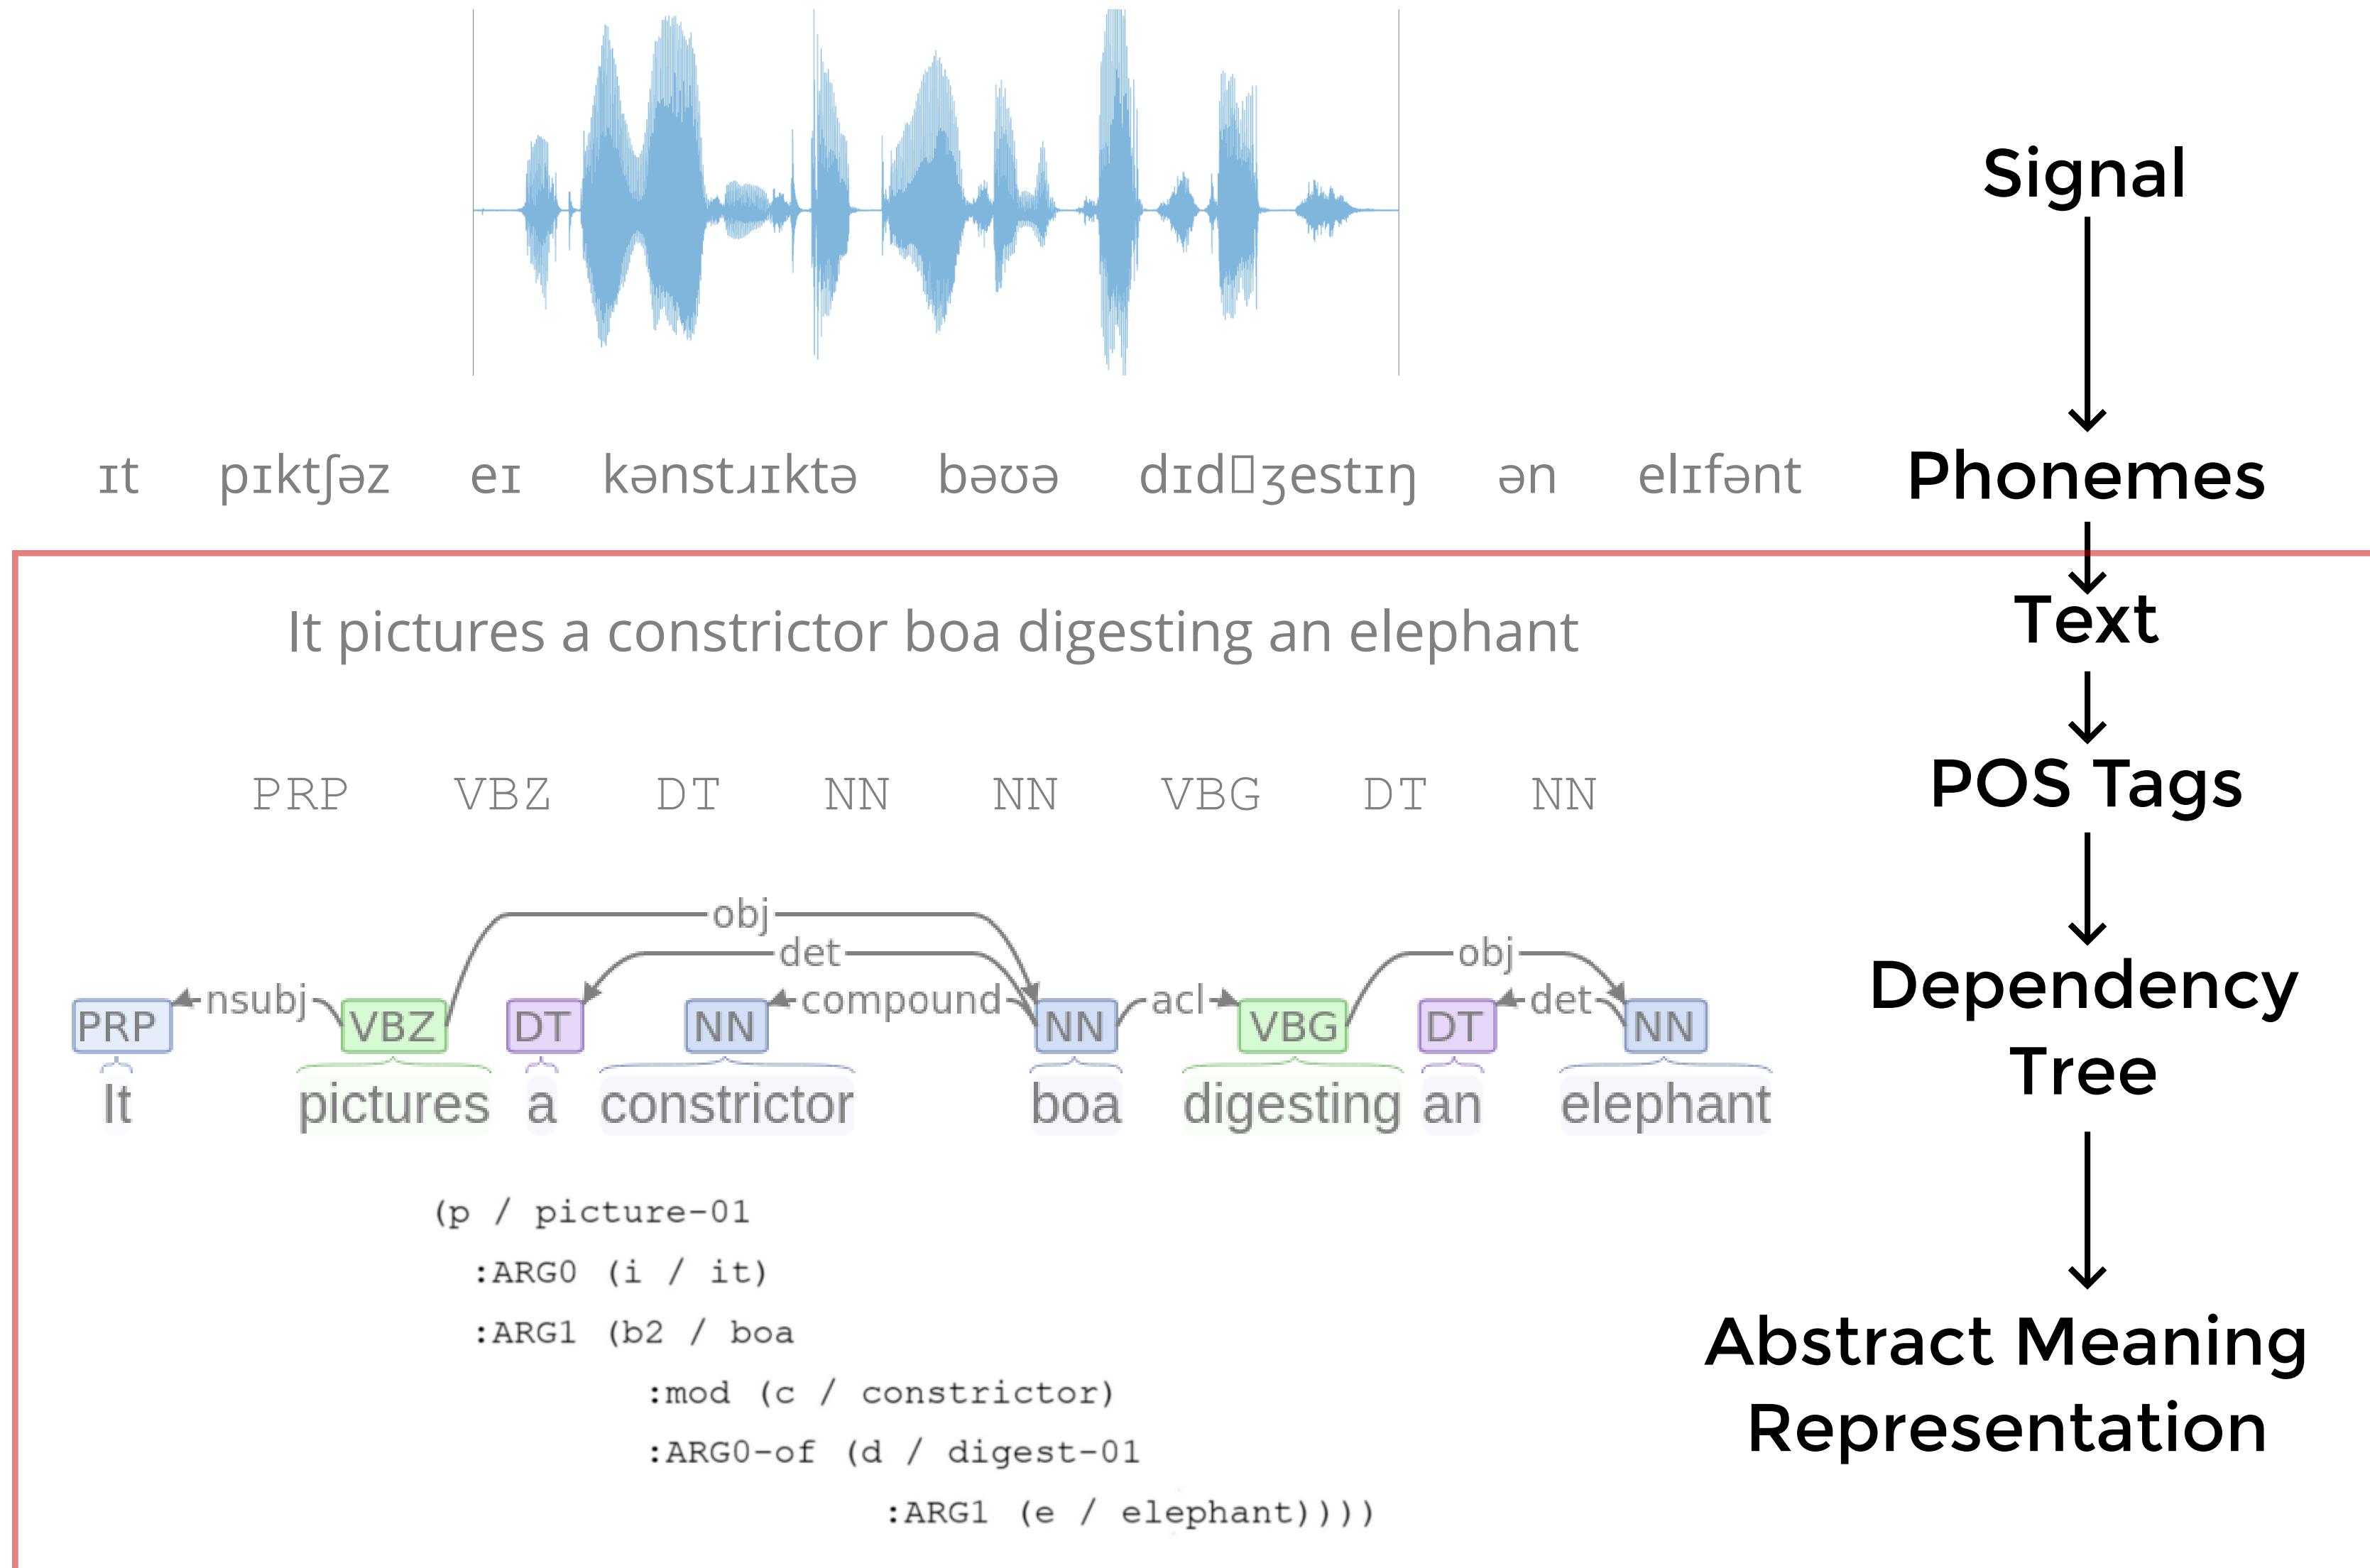
# Analyzing language: tasks involved



# Analyzing language: tasks involved



# Analyzing language: tasks involved



# Different kinds of linguistic knowledge

---

**To perform those tasks, what do we need ?**

# Different kinds of linguistic knowledge

---

**To perform those tasks, what do we need ?**

To enable completely understanding a statement, we need to work at different *levels*:

# Different kinds of linguistic knowledge

---

## To perform those tasks, what do we need ?

To enable completely understanding a statement, we need to work at different *levels*:

### Segmentation

- Segment text into lexical units (words)
- Not trivial: many possible uses for any punctuation symbol
- No typographic normalization; new uses (emoticons)

# Different kinds of linguistic knowledge

---

## To perform those tasks, what do we need ?

To enable completely understanding a statement, we need to work at different *levels*:

### Segmentation

- Segment text into lexical units (words)
- Not trivial: many possible uses for any punctuation symbol
- No typographic normalization; new uses (emoticons)

→ This is usually called **tokenization**

- Pieces are called *tokens*: not necessarily words anymore
- More on that later !

# Different kinds of linguistic knowledge

---

## Lexical treatment

- Identify words (*string to linguistic unit*) and their properties

# Different kinds of linguistic knowledge

---

## Lexical treatment

- Identify words (*string to linguistic unit*) and their properties
- Normalizing text and dealing with new words:
  - With the help of **morphology**
  - How words are built from *morphemes*

# Different kinds of linguistic knowledge

---

## Lexical treatment

- Identify words (*string to linguistic unit*) and their properties
  - Normalizing text and dealing with new words:
    - With the help of **morphology**
    - How words are built from *morphemes*
      - *Inflectional morphology*  
(*dog* → *dogs*, *play* → *played*)
- Variation:** same word, but modifies tense, number, ...

# Different kinds of linguistic knowledge

---

## Lexical treatment

- Identify words (*string to linguistic unit*) and their properties
- Normalizing text and dealing with new words:
  - With the help of **morphology**
  - How words are built from *morphemes*
    - *Inflectional morphology*  
(*dog* → *dogs*, *play* → *played*)  
**Variation:** same word, but modifies tense, number, ...
    - *Derivational morphology*  
(*happy* → *happiness*, *teach* → *teacher*)  
**Formation:** changes meaning, grammatical category

# Different kinds of linguistic knowledges

---

## Structural knowledge to assemble words: Syntax

- Constraints to obtain grammatically correct sentences: how words are organized  
→ Validity in *position* and *agreement*

# Different kinds of linguistic knowledges

## Structural knowledge to assemble words: Syntax

- Constraints to obtain grammatically correct sentences: how words are organized  
→ Validity in *position* and *agreement*
- Can be encoded in various ways (*constituency trees*, *dependency trees*) with their own pros and cons

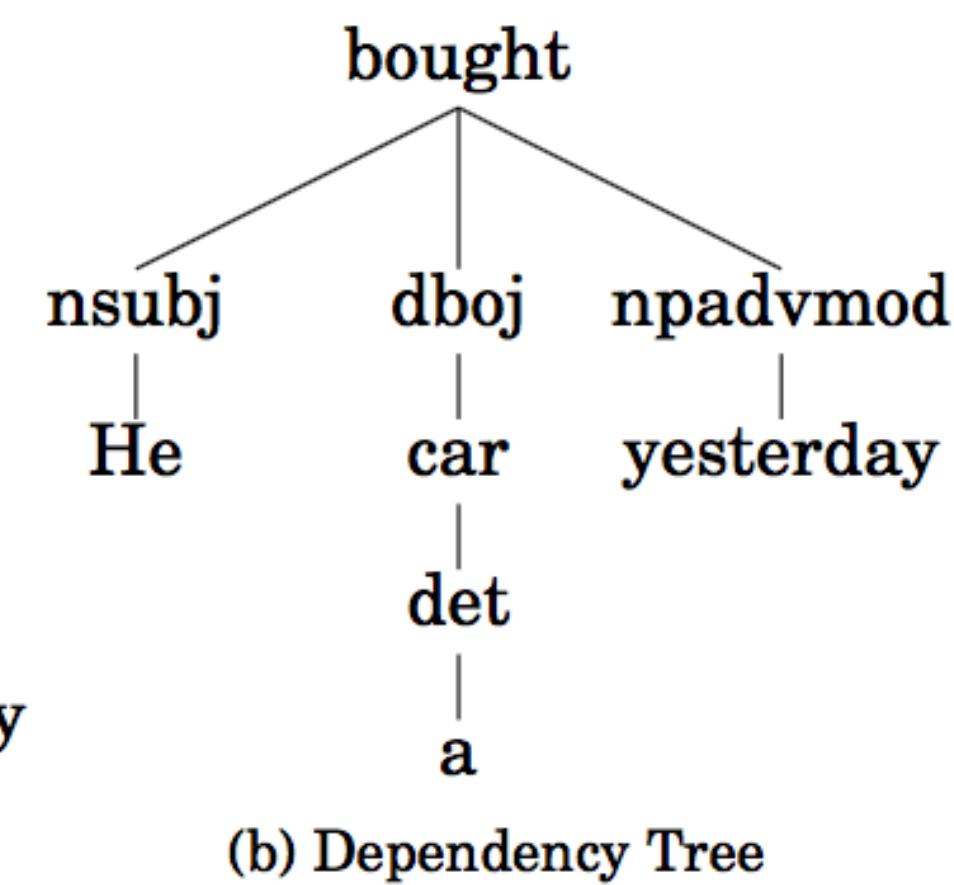
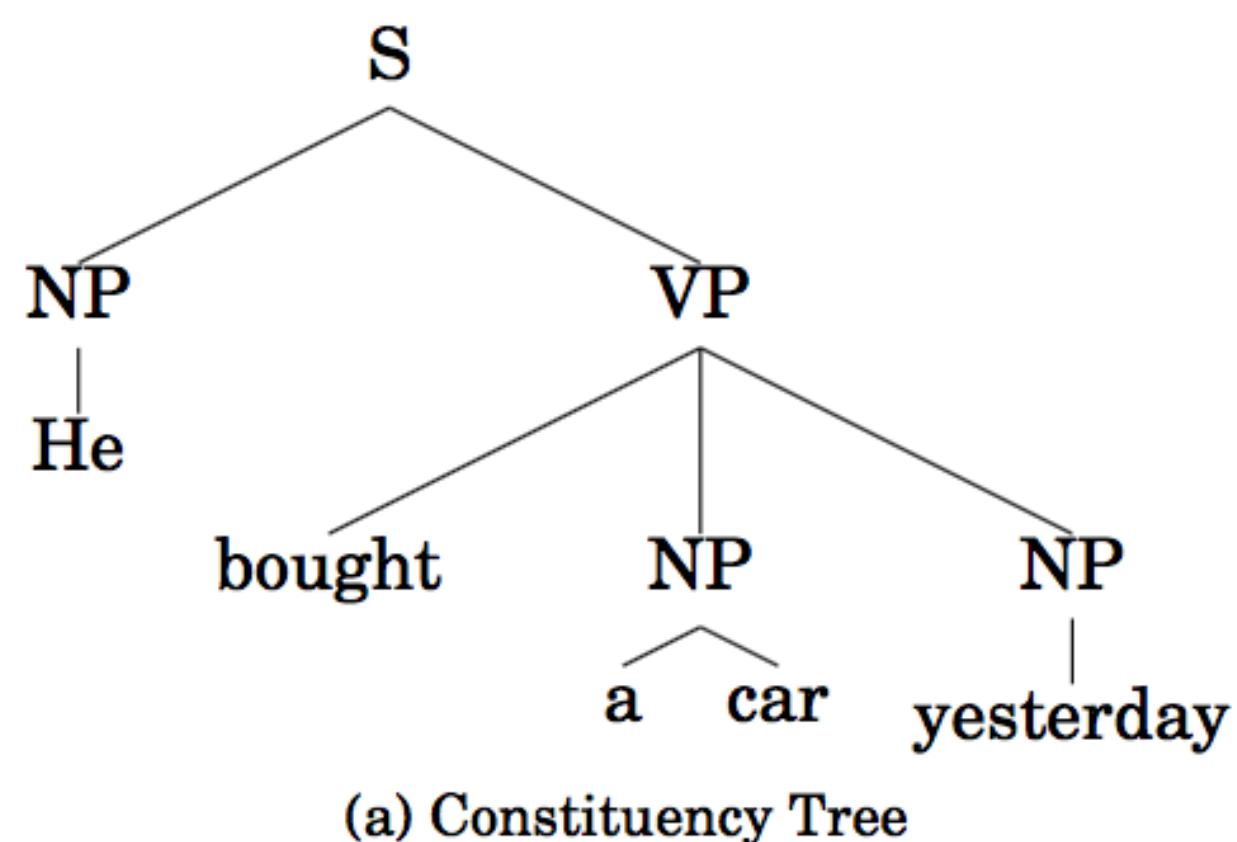


Fig. 3: Examples of the results of constituency and dependency parsing.

# Different kinds of linguistic knowledges

---

## Semantics

- How words and sentences carry a meaning
- Lexical semantics: meaning relationships between words  
(*synonymy, antonymy, hyponymy*)

# Different kinds of linguistic knowledges

---

## Semantics

- How words and sentences carry a meaning
- Lexical semantics: meaning relationships between words (*synonymy, antonymy, hyponymy*)
- **Compositionality**
  - Dealing with idioms: *kick the bucket* / Polysemy: *apple*

# Different kinds of linguistic knowledges

---

## Semantics

- How words and sentences carry a meaning
- Lexical semantics: meaning relationships between words (*synonymy, antonymy, hyponymy*)
- **Compositionality**
  - Dealing with idioms: *kick the bucket* / Polysemy: *apple*

## Pragmatics

- How context influences meaning
- Meaning depends on speaker intent, situation, and shared knowledge (implicit, social context)  
*Can you close the window ? / It's cold in here... / Oh, great, air conditionning in winter !*

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

- Ambiguity in written representation (speech synthesis)

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

- Ambiguity in written representation (speech synthesis)
- **Lexical ambiguity**, on word grammatical properties (Part-of-speech tagging) or sense (Word sense disambiguation)

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

- Ambiguity in written representation (speech synthesis)
- **Lexical ambiguity**, on word grammatical properties (Part-of-speech tagging) or sense (Word sense disambiguation)
- **Syntactical ambiguity** (parsing)

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

- Ambiguity in written representation (speech synthesis)
- **Lexical ambiguity**, on word grammatical properties (Part-of-speech tagging) or sense (Word sense disambiguation)
- **Syntactical ambiguity** (parsing)
- Ambiguity in **interpreting** a statement (sentiment classification, natural language inference)

# Why is NLP challenging ?

---

Most tasks in Natural Language Processing can be viewed as resolving *ambiguity* at one of different levels:

Learning to resolve natural language ambiguities: a unified approach (Dan Roth, 1998)

- Ambiguity in written representation (speech synthesis)
- **Lexical ambiguity**, on word grammatical properties (Part-of-speech tagging) or sense (Word sense disambiguation)
- **Syntactical ambiguity** (parsing)
- Ambiguity in **interpreting** a statement (sentiment classification, natural language inference)

Plus, *implicit* knowledge:

- Background, **commonsense knowledge**
- Contextual knowledge

makes things difficult...

# Winograd Schema Challenge

---

A *Winograd Schema* is a **small reading comprehension test** involving a single binary question

*The Winograd Schema Challenge (Levesque et al, 2012)*

# Winograd Schema Challenge

A *Winograd Schema* is a **small reading comprehension test** involving a single binary question

*The Winograd Schema Challenge (Levesque et al, 2012)*

The man couldn't lift his son because he was so **weak**. ————— Who was weak?

The man couldn't lift his son because he was so **heavy**. ————— Who was heavy?

Mary and Sue are **sisters**. } Mary and Sue are **mothers**. } How are Mary and Sue related?

Joan made sure to thank Susan for all the help she had **received**. ————— Who had received help?

Joan made sure to thank Susan for all the help she had **given**. ————— Who had given help?

John **promised** Bill to leave, so an hour later he left. } John **ordered** Bill to leave, so an hour later he left. } Who left an hour later?

Examples of ambiguity in language from the Winograd Schema Challenge

From "Practical Language Processing, Figure 1-7, Chapter 1"

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*
- Deep study of restricted areas: *natural language understanding* with Winograd, frameworks for *discourse modeling*

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*
- Deep study of restricted areas: *natural language understanding* with Winograd, frameworks for *discourse modeling*
- Probabilistic models used early for tasks like *optical character recognition* regained popularity towards the end of the 80s, with models like Hidden Markov Models (HMMs) for *speech recognition*

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*
- Deep study of restricted areas: *natural language understanding* with Winograd, frameworks for *discourse modeling*
- Probabilistic models used early for tasks like *optical character recognition* regained popularity towards the end of the 80s, with models like Hidden Markov Models (HMMs) for *speech recognition*
- Probabilistic **data-driven** models then rapidly became standard

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*
- Deep study of restricted areas: *natural language understanding* with Winograd, frameworks for *discourse modeling*
- Probabilistic models used early for tasks like *optical character recognition* regained popularity towards the end of the 80s, with models like Hidden Markov Models (HMMs) for *speech recognition*
- Probabilistic **data-driven** models then rapidly became standard
- Statistical models took over rapidly from 2000, thanks to:
  - A large amount of material and resources (+computational)
  - Efficiency of statistical learning (+unsupervised approaches)
  - Community effort: shared tasks, evaluation campaigns

# Statistical NLP

---

Historically, two paradigms: **symbolic** and **stochastic**

- On the symbolic side, formal systems based on logic and *grammars*
- Deep study of restricted areas: *natural language understanding* with Winograd, frameworks for *discourse modeling*
- Probabilistic models used early for tasks like *optical character recognition* regained popularity towards the end of the 80s, with models like Hidden Markov Models (HMMs) for *speech recognition*
- Probabilistic **data-driven** models then rapidly became standard
- Statistical models took over rapidly from 2000, thanks to:
  - A large amount of material and resources (+computational)
  - Efficiency of statistical learning (+unsupervised approaches)
  - Community effort: shared tasks, evaluation campaigns
- .. and now soon represented the state-of-the-art for almost any task; now, **deep learning**

# Schedule and instructions

---

- 12/09 - Introduction to **Text Processing and Symbolic Text Representations**
  - 19/09 - Text pre-processing, representations and visualization
- 26/09 - Introduction to **Language Modeling and Text Generation**
  - 03/10 - Introduction to language modeling
- 10/10 - **Word Embeddings**, Algorithms and Applications
  - 17/10 - Topic modeling and classification (Graded)
- 24/10 - **Sequence models**, Encoders and Decoders, Contextual Representations and Transfer Learning for NLP Tasks
  - 07/11 - Machine Translation (Graded, Mini project)
- 14/11 - **Structured Prediction** in NLP
  - 21/11 - Structured prediction

# Schedule and instructions

---

- 28/11 - **Large Language Models and Societal Impact**
- 05/12 - **Sentiment Analysis: Towards Transparent Models**
- 12/12 - **Conversational Systems**

# Schedule and instructions

---

- 28/11 - **Large Language Models and Societal Impact**
- 05/12 - **Sentiment Analysis: Towards Transparent Models**
- 12/12 - **Conversational Systems**

## Evaluation:

- **Labs 3** (10%) and **4** (20%) will be graded
- **Exam** (70%, with limited personal notes)
  - Questions focused on the course: *general knowledge, methodology, good practices*
  - Small exercises *based on the labs*

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

Speech and Language  
Processing (Jurafsky &  
Manning) (Chapter II)

User: I am unhappy.  
ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY  
User: I need some help, that much seems certain.  
ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP  
User: Perhaps I could learn to get along with my mother.  
ELIZA: TELL ME MORE ABOUT YOUR FAMILY  
User: My mother takes care of me.  
ELIZA: WHO ELSE IN YOU FAMILY TAKES CARE OF YOU  
User: My father.  
ELIZA: YOUR FATHER  
User: You are like my father in some ways.

Weizenbaum (1966)

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

Speech and Language  
Processing (Jurafsky &  
Manning) (Chapter II)

User: I am unhappy.  
ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY  
User: I need some help, that much seems certain.  
ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP  
User: Perhaps I could learn to get along with my mother.  
ELIZA: TELL ME MORE ABOUT YOUR FAMILY  
User: My mother takes care of me.  
ELIZA: WHO ELSE IN YOU FAMILY TAKES CARE OF YOU  
User: My father.  
ELIZA: YOUR FATHER  
User: You are like my father in some ways.

Weizenbaum (1966)

- Uses **pattern matching** to recognize phrases

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

Speech and Language  
Processing (Jurafsky &  
Manning) (Chapter II)

User: I am unhappy.  
ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY  
User: I need some help, that much seems certain.  
ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP  
User: Perhaps I could learn to get along with my mother.  
ELIZA: TELL ME MORE ABOUT YOUR FAMILY  
User: My mother takes care of me.  
ELIZA: WHO ELSE IN YOU FAMILY TAKES CARE OF YOU  
User: My father.  
ELIZA: YOUR FATHER  
User: You are like my father in some ways.

Weizenbaum (1966)

- Uses **pattern matching** to recognize phrases
- Uses a set of **rules** to translate them into suitable outputs

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

Speech and Language  
Processing (Jurafsky &  
Manning) (Chapter II)

User: I am unhappy.  
ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY  
User: I need some help, that much seems certain.  
ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP  
User: Perhaps I could learn to get along with my mother.  
ELIZA: TELL ME MORE ABOUT YOUR FAMILY  
User: My mother takes care of me.  
ELIZA: WHO ELSE IN YOUR FAMILY TAKES CARE OF YOU  
User: My father.  
ELIZA: YOUR FATHER  
User: You are like my father in some ways.

Weizenbaum (1966)

- Uses **pattern matching** to recognize phrases
- Uses a set of **rules** to translate them into suitable outputs
- Remarkably successful !

# Basic text processing

---

What is the bare minimum to process text ?

→ Example of **ELIZA**

Speech and Language  
Processing (Jurafsky &  
Manning) (Chapter II)

User: I am unhappy.  
ELIZA: DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY  
User: I need some help, that much seems certain.  
ELIZA: WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP  
User: Perhaps I could learn to get along with my mother.  
ELIZA: TELL ME MORE ABOUT YOUR FAMILY  
User: My mother takes care of me.  
ELIZA: WHO ELSE IN YOUR FAMILY TAKES CARE OF YOU  
User: My father.  
ELIZA: YOUR FATHER  
User: You are like my father in some ways.

Weizenbaum (1966)

- Uses **pattern matching** to recognize phrases
  - Uses a set of **rules** to translate them into suitable outputs
  - Remarkably successful !
- For pattern matching: **regular expressions** !
- Used then for **text normalization** and **tokenization**

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for '*the*':

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences:

the

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences:
- Will find undesired ones:

the  
[tT]he

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences: the
- Will find undesired ones: [tT] he
- Will probably have very few false positives or negatives: [ ^a-zA-Z ] [tT] he [ ^a-zA-Z ]

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences: the
- Will find undesired ones: [tT] he
- Will probably have very few false positives or negatives: [ ^a-zA-Z ] [tT] he [ ^a-zA-Z ]

→ Can be used to capture and substitute text !

- Replacing 'the' with 'The' s/the/The

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences: the
- Will find undesired ones: [tT] he
- Will probably have very few false positives or negatives: [ ^a-zA-Z ] [tT] he [ ^a-zA-Z ]

→ Can be used to capture and substitute text !

- Replacing 'the' with 'The' s/the/The
- Capturing any string ending with 'er': / (.\*) er/

# Regular expressions

---

An algebraic notation for characterizing a set of strings - used practically everywhere !

→ Can be used to find desired occurrences of words in a large text !

Example: looking for 'the':

- Will miss some occurrences: the
- Will find undesired ones: [tT] he
- Will probably have very few false positives or negatives: [ ^a-zA-Z ] [tT] he [ ^a-zA-Z ]

→ Can be used to capture and substitute text !

- Replacing 'the' with 'The' s/the/The
- Capturing any string ending with 'er': / (.\*)er/
- Getting a superlative: s/the (.\*)er/the (\1)est/

# Regular expressions and ELIZA

---

Regular expressions also allows for more complex functionalities.. but especially, allows for **ELIZA**:

- Early NLP system that imitated a *Rogerian psychotherapist*, by Joseph Weizenbaum (1966)
- Has a series or cascade of regular expression substitutions, matching and changing some part of the input lines:

# Regular expressions and ELIZA

---

Regular expressions also allows for more complex functionalities.. but especially, allows for **ELIZA**:

- Early NLP system that imitated a *Rogerian psychotherapist*, by Joseph Weizenbaum (1966)
- Has a series or cascade of regular expression substitutions, matching and changing some part of the input lines:
  - Input lines are uppercased
  - Change all instances of MY to YOUR, I'M to YOU ARE, etc...
  - Then, other patterns are replaced:

# Regular expressions and ELIZA

---

Regular expressions also allows for more complex functionalities.. but especially, allows for **ELIZA**:

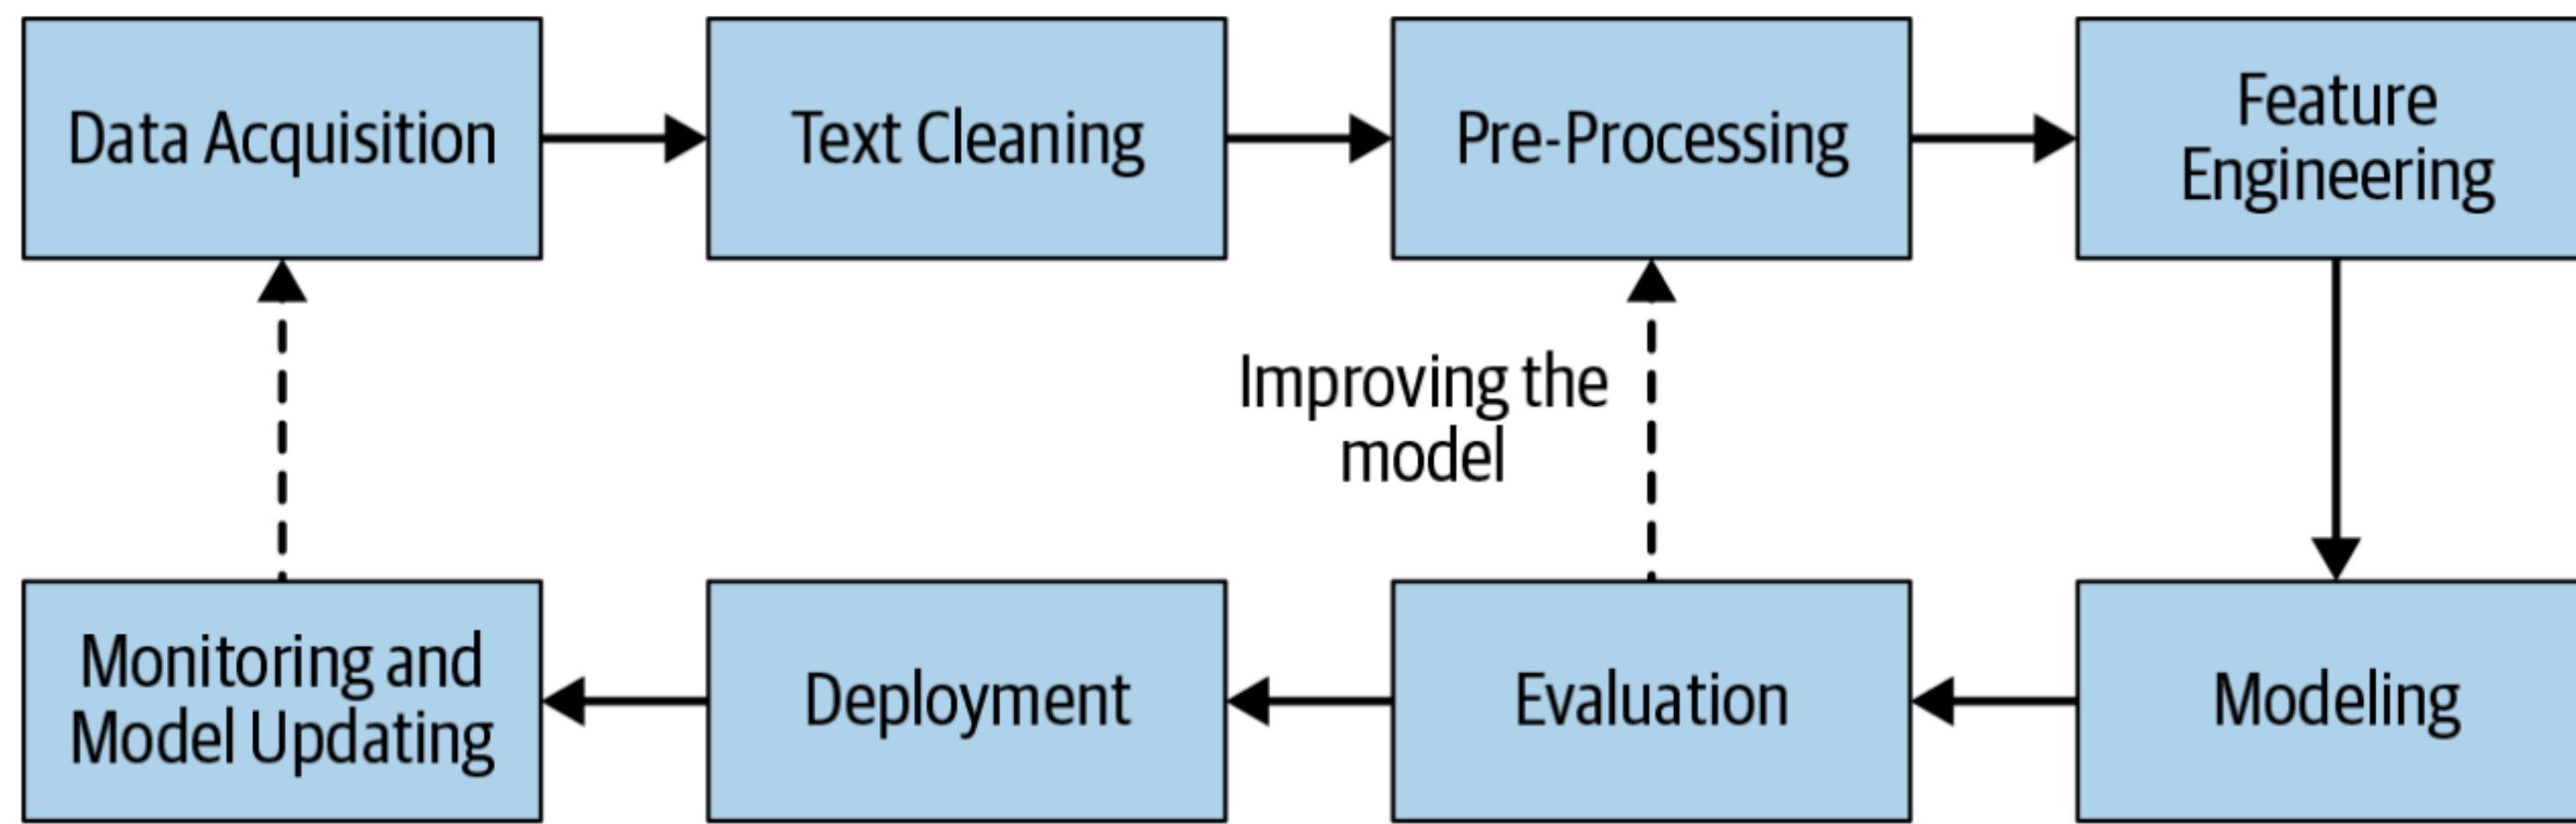
→ Early NLP system that imitated a *Rogerian psychotherapist*, by Joseph Weizenbaum (1966)

→ Has a series or cascade of regular expression substitutions, matching and changing some part of the input lines:

- Input lines are uppercased
- Change all instances of MY to YOUR, I'M to YOU ARE, etc...
- Then, other patterns are replaced:

```
s/.* I'M (depressed|sad) .*/I AM SORRY TO HEAR YOU ARE \1/  
s/.* I AM (depressed|sad) .*/WHY DO YOU THINK YOU ARE \1/  
          s/.* all .*/IN WHAT WAY/  
s/.* always .*/CAN YOU THINK OF A SPECIFIC EXAMPLE/
```

# How to build an NLP application ?



General NLP Pipeline. *From "Practical Language Processing, Figure 2-1, Chapter 2"*

## Crucial steps:

- **Acquiring data** - even for rule-based system, if only for *evaluation*
- **Pre-processing**
- Feature Engineering → classical approaches *vs* modern **Text representation** through deep learning
- **Evaluation** - *intrinsic* and *extrinsic*

# A few definitions

---

What are the *objects* we will work on ?

# A few definitions

---

What are the *objects* we will work on ?

- We begin from a **corpus** (plural **corpora**): computer-readable collection of text, broken down into *words*

# A few definitions

---

What are the *objects* we will work on ?

- We begin from a **corpus** (plural **corpora**): computer-readable collection of text, broken down into *words*
- **Word types** are the number of distinct words in a corpus - usually grouped in a set, the **vocabulary**

# A few definitions

---

What are the *objects* we will work on ?

- We begin from a **corpus** (plural **corpora**): computer-readable collection of text, broken down into *words*
- **Word types** are the number of distinct words in a corpus - usually grouped in a set, the **vocabulary**
- **Word tokens** are instances of types in the running text

# A few definitions

---

What are the *objects* we will work on ?

- We begin from a **corpus** (plural **corpora**): computer-readable collection of text, broken down into *words*
- **Word types** are the number of distinct words in a corpus - usually grouped in a set, the **vocabulary**
- **Word tokens** are instances of types in the running text
- Two types can be different **word forms** of a same **lemma**:  
→ *cat* and *cats* have the same lemma *cat*

# A few definitions

---

What are the *objects* we will work on ?

- We begin from a **corpus** (plural **corpora**): computer-readable collection of text, broken down into *words*
- **Word types** are the number of distinct words in a corpus - usually grouped in a set, the **vocabulary**
- **Word tokens** are instances of types in the running text
- Two types can be different **word forms** of a same **lemma**:  
→ *cat* and *cats* have the same lemma *cat*

→ Example:

*"I showed my masterpiece to the grown-ups, and asked them whether the drawing frightened them."*

# Acquiring data: corpora and resources

---

Data-driven methods are based on **corpora**, which have many distinct properties:

# Acquiring data: corpora and resources

---

Data-driven methods are based on **corpora**, which have many distinct properties:

- Language (7000+) and varieties, code switching...
- Genre (news, scientific, fiction..), specific domain (medical, law...)
- Source and authors: how was it written ? Collected ? Why ?
- Use *data statement* to avoid biases

*Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science (Bender & Friedman, 2018)*

# Acquiring data: corpora and resources

---

Data-driven methods are based on **corpora**, which have many distinct properties:

- Language (7000+) and varieties, code switching...
- Genre (news, scientific, fiction..), specific domain (medical, law...)
- Source and authors: how was it written ? Collected ? Why ?
- Use *data statement* to avoid biases

*Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science (Bender & Friedman, 2018)*

Resources are unevenly distributed along languages ! Still, they are very diverse. For example:

- Lexical Databases like **WordNet** (but also for other languages)
- Labeled data for many tasks - will allow supervised learning
- Careful: *annotation* is a difficult and subjective process

# Pre-processing: Tokenization

---

The first step to any task: pre-process the text, which begins by segmenting it into words. This is **tokenization** !

# Pre-processing: Tokenization

---

The first step to any task: pre-process the text, which begins by segmenting it into words. This is **tokenization** !

- Simplest approach: *space-based* - segments along spaces  
→ Does not work with some languages !

# Pre-processing: Tokenization

---

The first step to any task: pre-process the text, which begins by segmenting it into words. This is **tokenization** !

- Simplest approach: *space-based* - segments along spaces  
→ Does not work with some languages !
- What to do with punctuation ? Example:

*'But they answered: "Frighten? Why should any one be frightened by a hat?" My drawing was not a picture of a hat.'*

and many other depending on the context (hashtags, emails, etc...)

# Pre-processing: Tokenization

---

The first step to any task: pre-process the text, which begins by segmenting it into words. This is **tokenization** !

- Simplest approach: *space-based* - segments along spaces  
→ Does not work with some languages !
- What to do with punctuation ? Example:

*'But they answered: "Frighten? Why should any one be frightened by a hat?" My drawing was not a picture of a hat.'*

- and many other depending on the context (hashtags, emails, etc...)
- Tools: we will use packages like **NLTK** and **Spacy**
  - Tokenizers are based on rules, with regular expression

# Pre-processing: Tokenization

---

The first step to any task: pre-process the text, which begins by segmenting it into words. This is **tokenization** !

- Simplest approach: *space-based* - segments along spaces  
→ Does not work with some languages !
- What to do with punctuation ? Example:

*'But they answered: "Frighten? Why should any one be frightened by a hat?" My drawing was not a picture of a hat.'*

- and many other depending on the context (hashtags, emails, etc...)
  - Tools: we will use packages like **NLTK** and **Spacy**
    - Tokenizers are based on rules, with regular expression
- After **tokenization**, the set of tokens are listed in a vocabulary  $\mathcal{V}$  (*word types*)

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

- What to do with upper-cases ?  
→ depends on the task ! Information retrieval does not need them  
- but information extraction does !

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

- What to do with upper-cases ?  
→ depends on the task ! Information retrieval does not need them  
- but information extraction does !
- **Lemmatization:** represent words as their lemma
  - Done by morphological parsing
  - Necessary for some languages: e.g, Turkish.
  - Tools - we can also use here **NLTK**

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

- What to do with upper-cases ?  
→ depends on the task ! Information retrieval does not need them  
- but information extraction does !
- **Lemmatization**: represent words as their lemma
  - Done by morphological parsing
  - Necessary for some languages: e.g, Turkish.
  - Tools - we can also use here **NLTK**
- Other process: **stemming** - crudely cutting words

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

- What to do with upper-cases ?  
→ depends on the task ! Information retrieval does not need them  
- but information extraction does !
- **Lemmatization**: represent words as their lemma
  - Done by morphological parsing
  - Necessary for some languages: e.g, Turkish.
  - Tools - we can also use here **NLTK**
- Other process: **stemming** - crudely cutting words

Stemming	Lemmatization
adjustable -> adjust	was -> (to) be
formality -> formaliti	better -> good
formaliti -> formal	meeting -> meeting
airliner -> airlin	

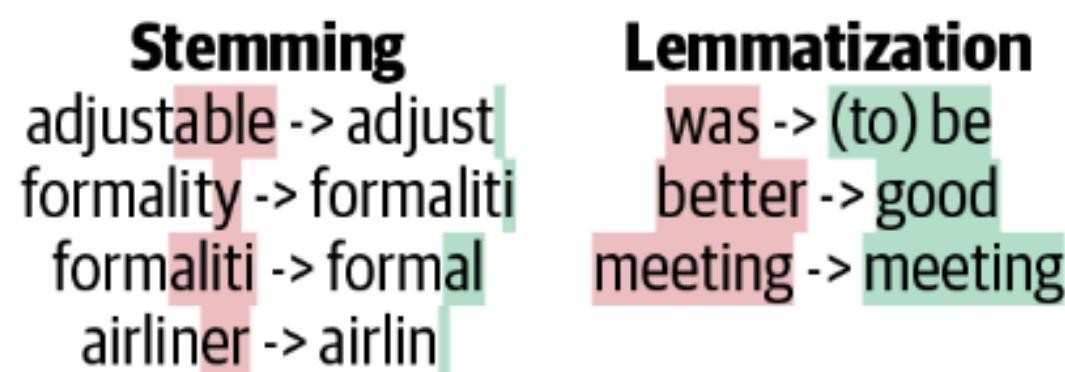
*From "Practical Language Processing,  
Figure 2-7, Chapter 2"*

# Pre-processing: Word Normalization

---

We may want to choose a standard format for words:

- What to do with upper-cases ?  
→ depends on the task ! Information retrieval does not need them  
- but information extraction does !
- **Lemmatization**: represent words as their lemma
  - Done by morphological parsing
  - Necessary for some languages: e.g, Turkish.
  - Tools - we can also use here **NLTK**
- Other process: **stemming** - crudely cutting words



*From "Practical Language Processing,  
Figure 2-7, Chapter 2"*

→ Let's think ahead: what would be the advantages of the various strategies for text normalization ?

# Issues with Tokenization

---

- **Word-level tokenization:**
  - Treats different forms of the same root as separate (e.g., “low”, “lowest”, “lowered”, etc)
  - This implies separate parameters !
  - Issues with **unknown words**

# Issues with Tokenization

---

- **Word-level tokenization:**
  - Treats different forms of the same root as separate (e.g., “low”, “lowest”, “lowered”, etc)
  - This implies separate parameters !
  - Issues with **unknown words**
- **Character-level Tokenization:**
  - Small vocabulary, no unknown, but far longer sequences
  - How to combine character representations into representations of words meaning **must be learned**

# Issues with Tokenization

---

- **Word-level tokenization:**
  - Treats different forms of the same root as separate (e.g., “low”, “lowest”, “lowered”, etc)
  - This implies separate parameters !
  - Issues with **unknown words**
- **Character-level Tokenization:**
  - Small vocabulary, no unknown, but far longer sequences
  - How to combine character representations into representations of words meaning **must be learned**
- A middle ground ?
  - **Byte Pair Encoding (1994)** used for **Subword Tokenization**  
*Neural machine translation of rare words with subword units (Sennrich et al, 2016)*
  - Now the norm for modern models

# Byte Pair Encoding: Algorithm

---

We usually begin from *pre-tokenized* word-level tokens:

- $\mathcal{V} \leftarrow$  All characters in the training data, including an *end of word character*
- While  $|\mathcal{V}| < K$ :
  - Tokenize the data with the current  $\mathcal{V}$ 
    - Taking the *longest possible prefix* each time
  - Count the frequency of adjacent token pairs in the data
  - Choose the pair  $\{l, r\}$  that occurs most frequently
  - Add the pair to the vocabulary as a new token:  $\mathcal{V} \leftarrow \mathcal{V} \cup \{lr\}$

The algorithm saves the **merges**, **in order**, and applies them in order when splitting new words

**Example:**

# Byte Pair Encoding: Algorithm

---

We usually begin from *pre-tokenized* word-level tokens:

- $\mathcal{V} \leftarrow$  All characters in the training data, including an *end of word character*
- While  $|\mathcal{V}| < K$ :
  - Tokenize the data with the current  $\mathcal{V}$ 
    - Taking the *longest possible prefix* each time
  - Count the frequency of adjacent token pairs in the data
  - Choose the pair  $\{l, r\}$  that occurs most frequently
  - Add the pair to the vocabulary as a new token:  $\mathcal{V} \leftarrow \mathcal{V} \cup \{lr\}$

The algorithm saves the **merges**, **in order**, and applies them in order when splitting new words

	l o w _ : 5
	l o w e r _ : 2
<b>Example:</b>	n e w e s t _ : 6
	w i d e s t _ : 3

# Properties and variants

---

- Usually include frequent words and frequent subwords
  - Often, **morphemes** (*the smallest meaning-bearing unit of a language*)

Supercalifragilisticexpialidocious

- **WordPiece** (*Schuster et al., 2012*):
  - Used by Google for their models
  - Merge is different (maximizes *likelihood*, not frequency)
  - Merge order is not saved: tokenization uses *the longest subword first*
- **SentencePiece** (*Kudo et al., 2018*): applied to any algorithm
  - Replaces *whitespaces* by a special token and does not apply pre-tokenization
  - Very useful for languages without segmentation between words
- Many other variants (using *bytes*, *soft tokenization*...)

# Sentence tokenization

---

## What is a sentence ?

The format will heavily depend on the target task

- For many classification tasks, labels are at a higher level
  - **Labeled document:** possibly many sentences
- For core NLP tasks, it's more complicated

# Sentence tokenization

---

## What is a sentence ?

The format will heavily depend on the target task

- For many classification tasks, labels are at a higher level
  - **Labeled document:** possibly many sentences
- For core NLP tasks, it's more complicated
  - Ideally, it:
    - has a **complete syntactic structure**
    - is **semantically related to other sentences**

# Sentence tokenization

---

## What is a sentence ?

The format will heavily depend on the target task

- For many classification tasks, labels are at a higher level
  - **Labeled document:** possibly many sentences
- For core NLP tasks, it's more complicated
  - Ideally, it:
    - has a **complete syntactic structure**
    - is **semantically related to other sentences**
  - In practice, it is *typographically marked*, but:
    - The full stop may be ambiguous, the uppercase too
    - Difficulty with embedded sentences (e.g, with quotes)
    - No clear markers in some languages

# Tokens versus type

---

**Type:** an element of the vocabulary

**Token:** an instance of that type in running text

Relationship ?

# Tokens versus type

---

**Type:** an element of the vocabulary

**Token:** an instance of that type in running text

Relationship ?

- $N$  = Number of tokens
- $|\mathcal{V}|$  = Number of types
- **Heaps Law = Herdan's Law:**

$$|\mathcal{V}| = kN^\beta$$

where  $k$  and  $\beta$  are determined empirically

# Tokens versus type

---

**Type:** an element of the vocabulary

**Token:** an instance of that type in running text

Relationship ?

- $N$  = Number of tokens
- $|\mathcal{V}|$  = Number of types
- **Heaps Law = Herdan's Law:**

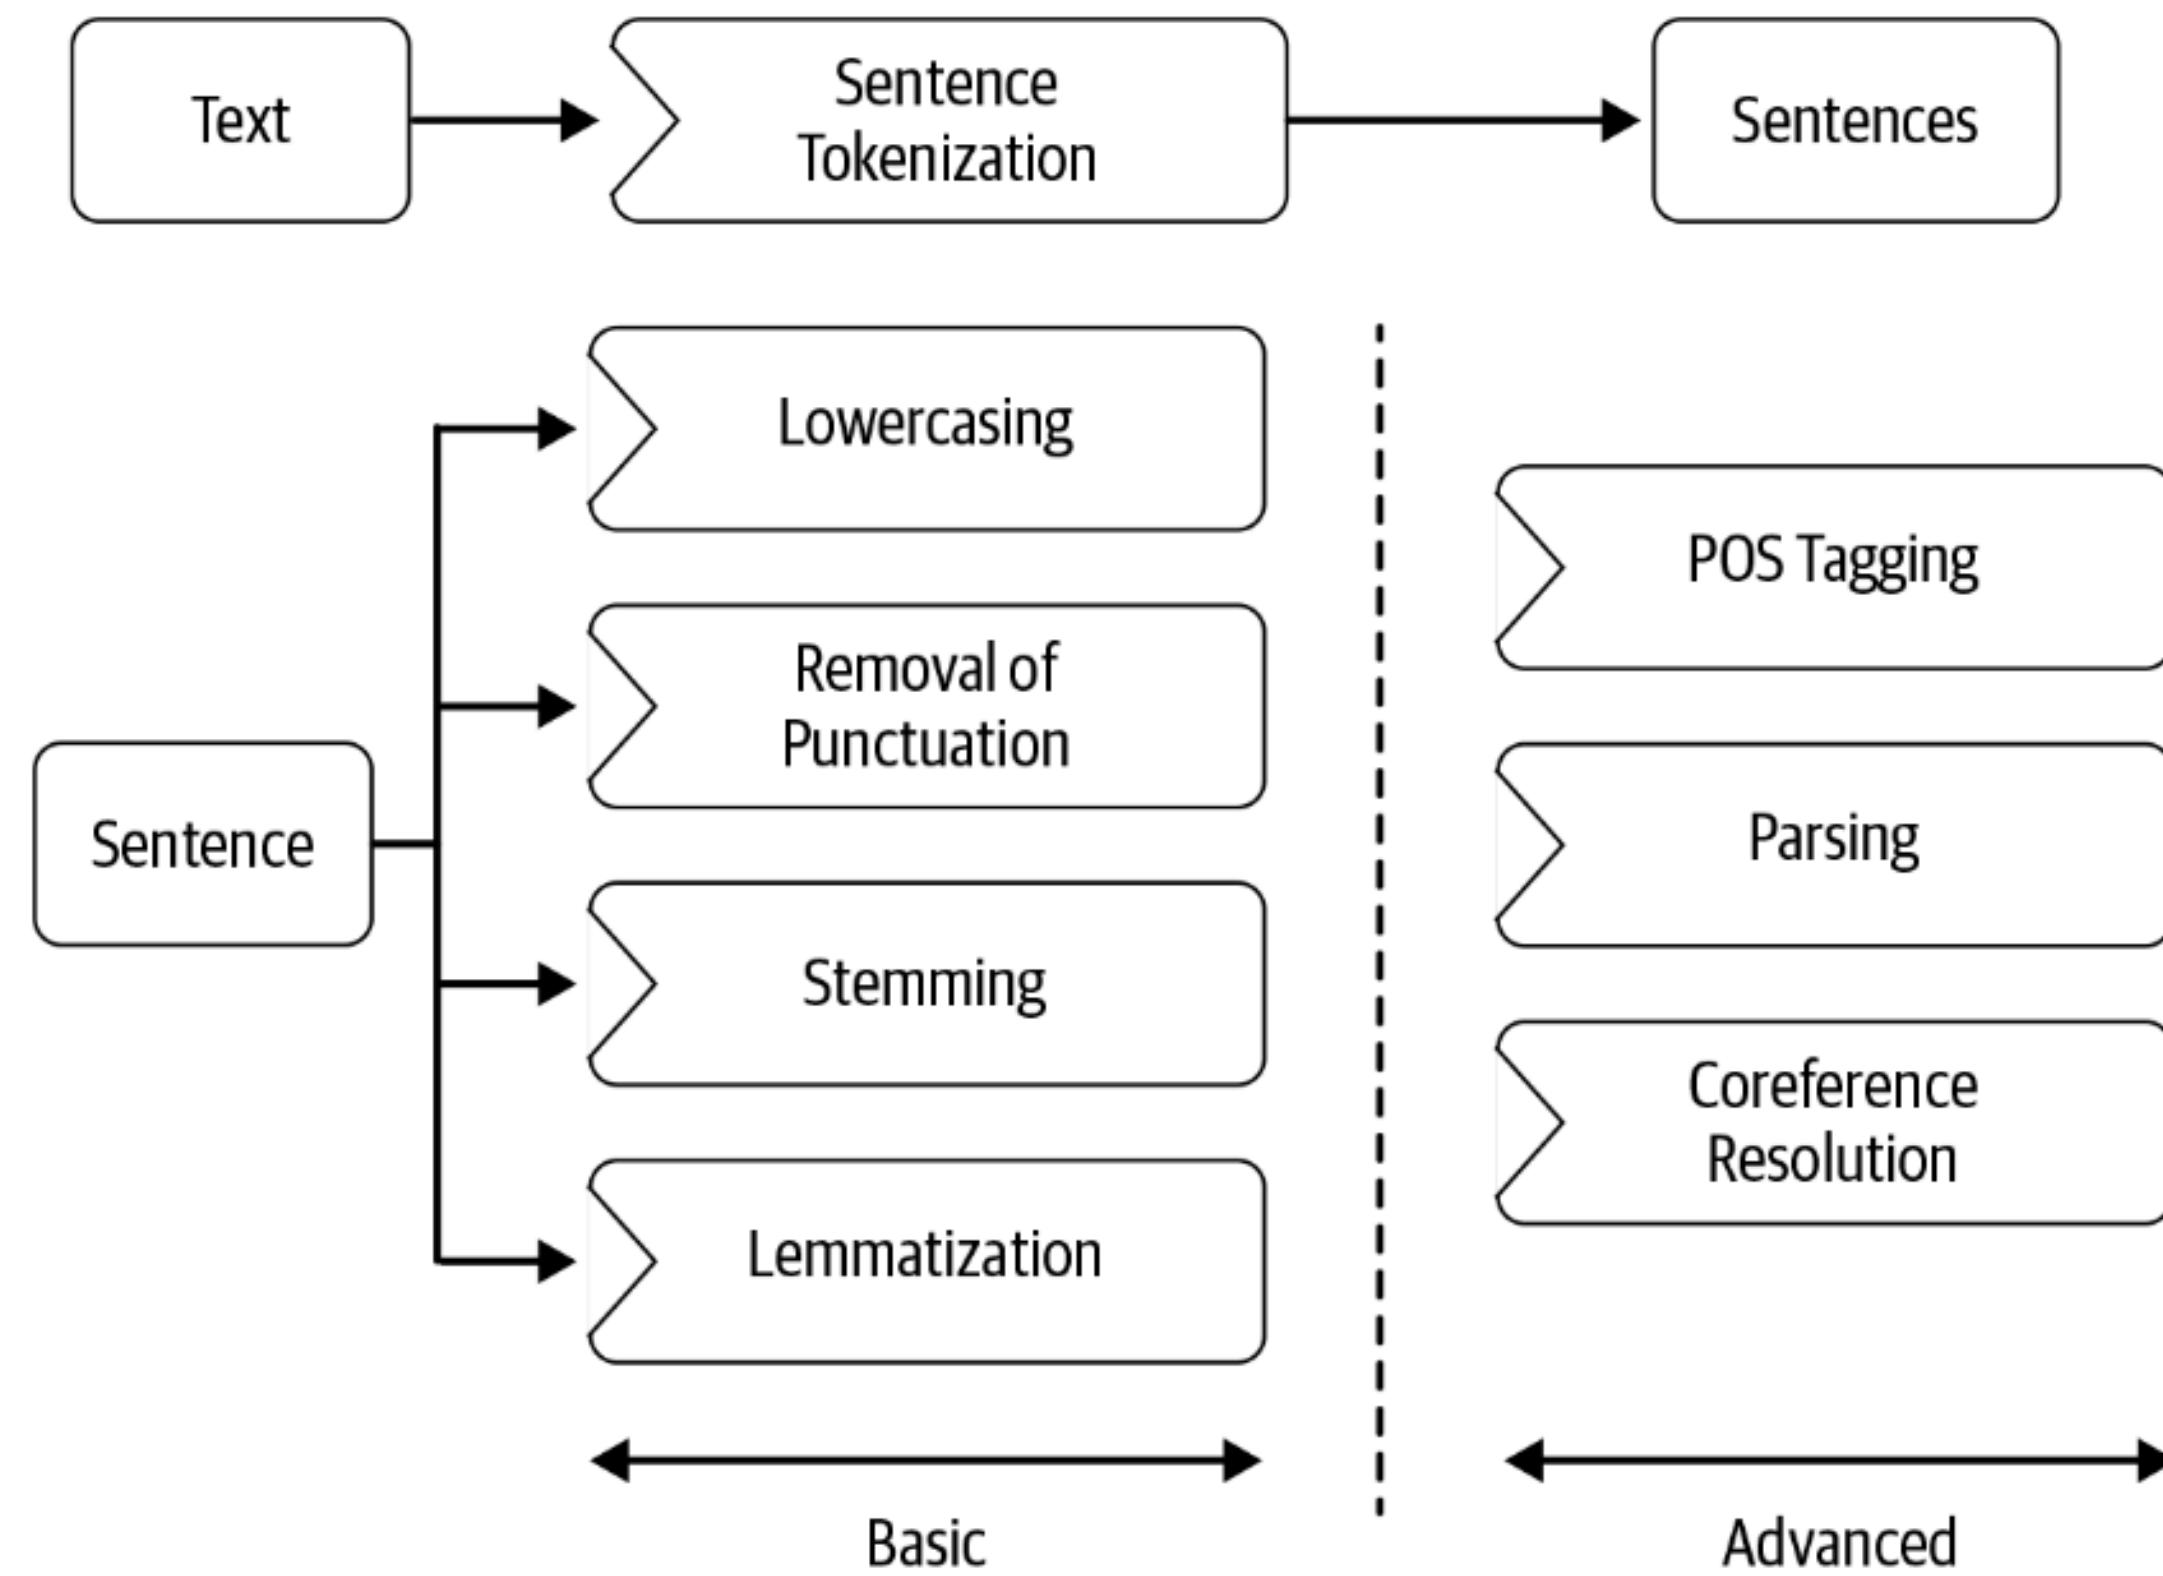
$$|\mathcal{V}| = kN^\beta$$

where  $k$  and  $\beta$  are determined empirically

- For English text corpora,  $10 < k < 100$  and  $0.4 < \beta < 0.6$

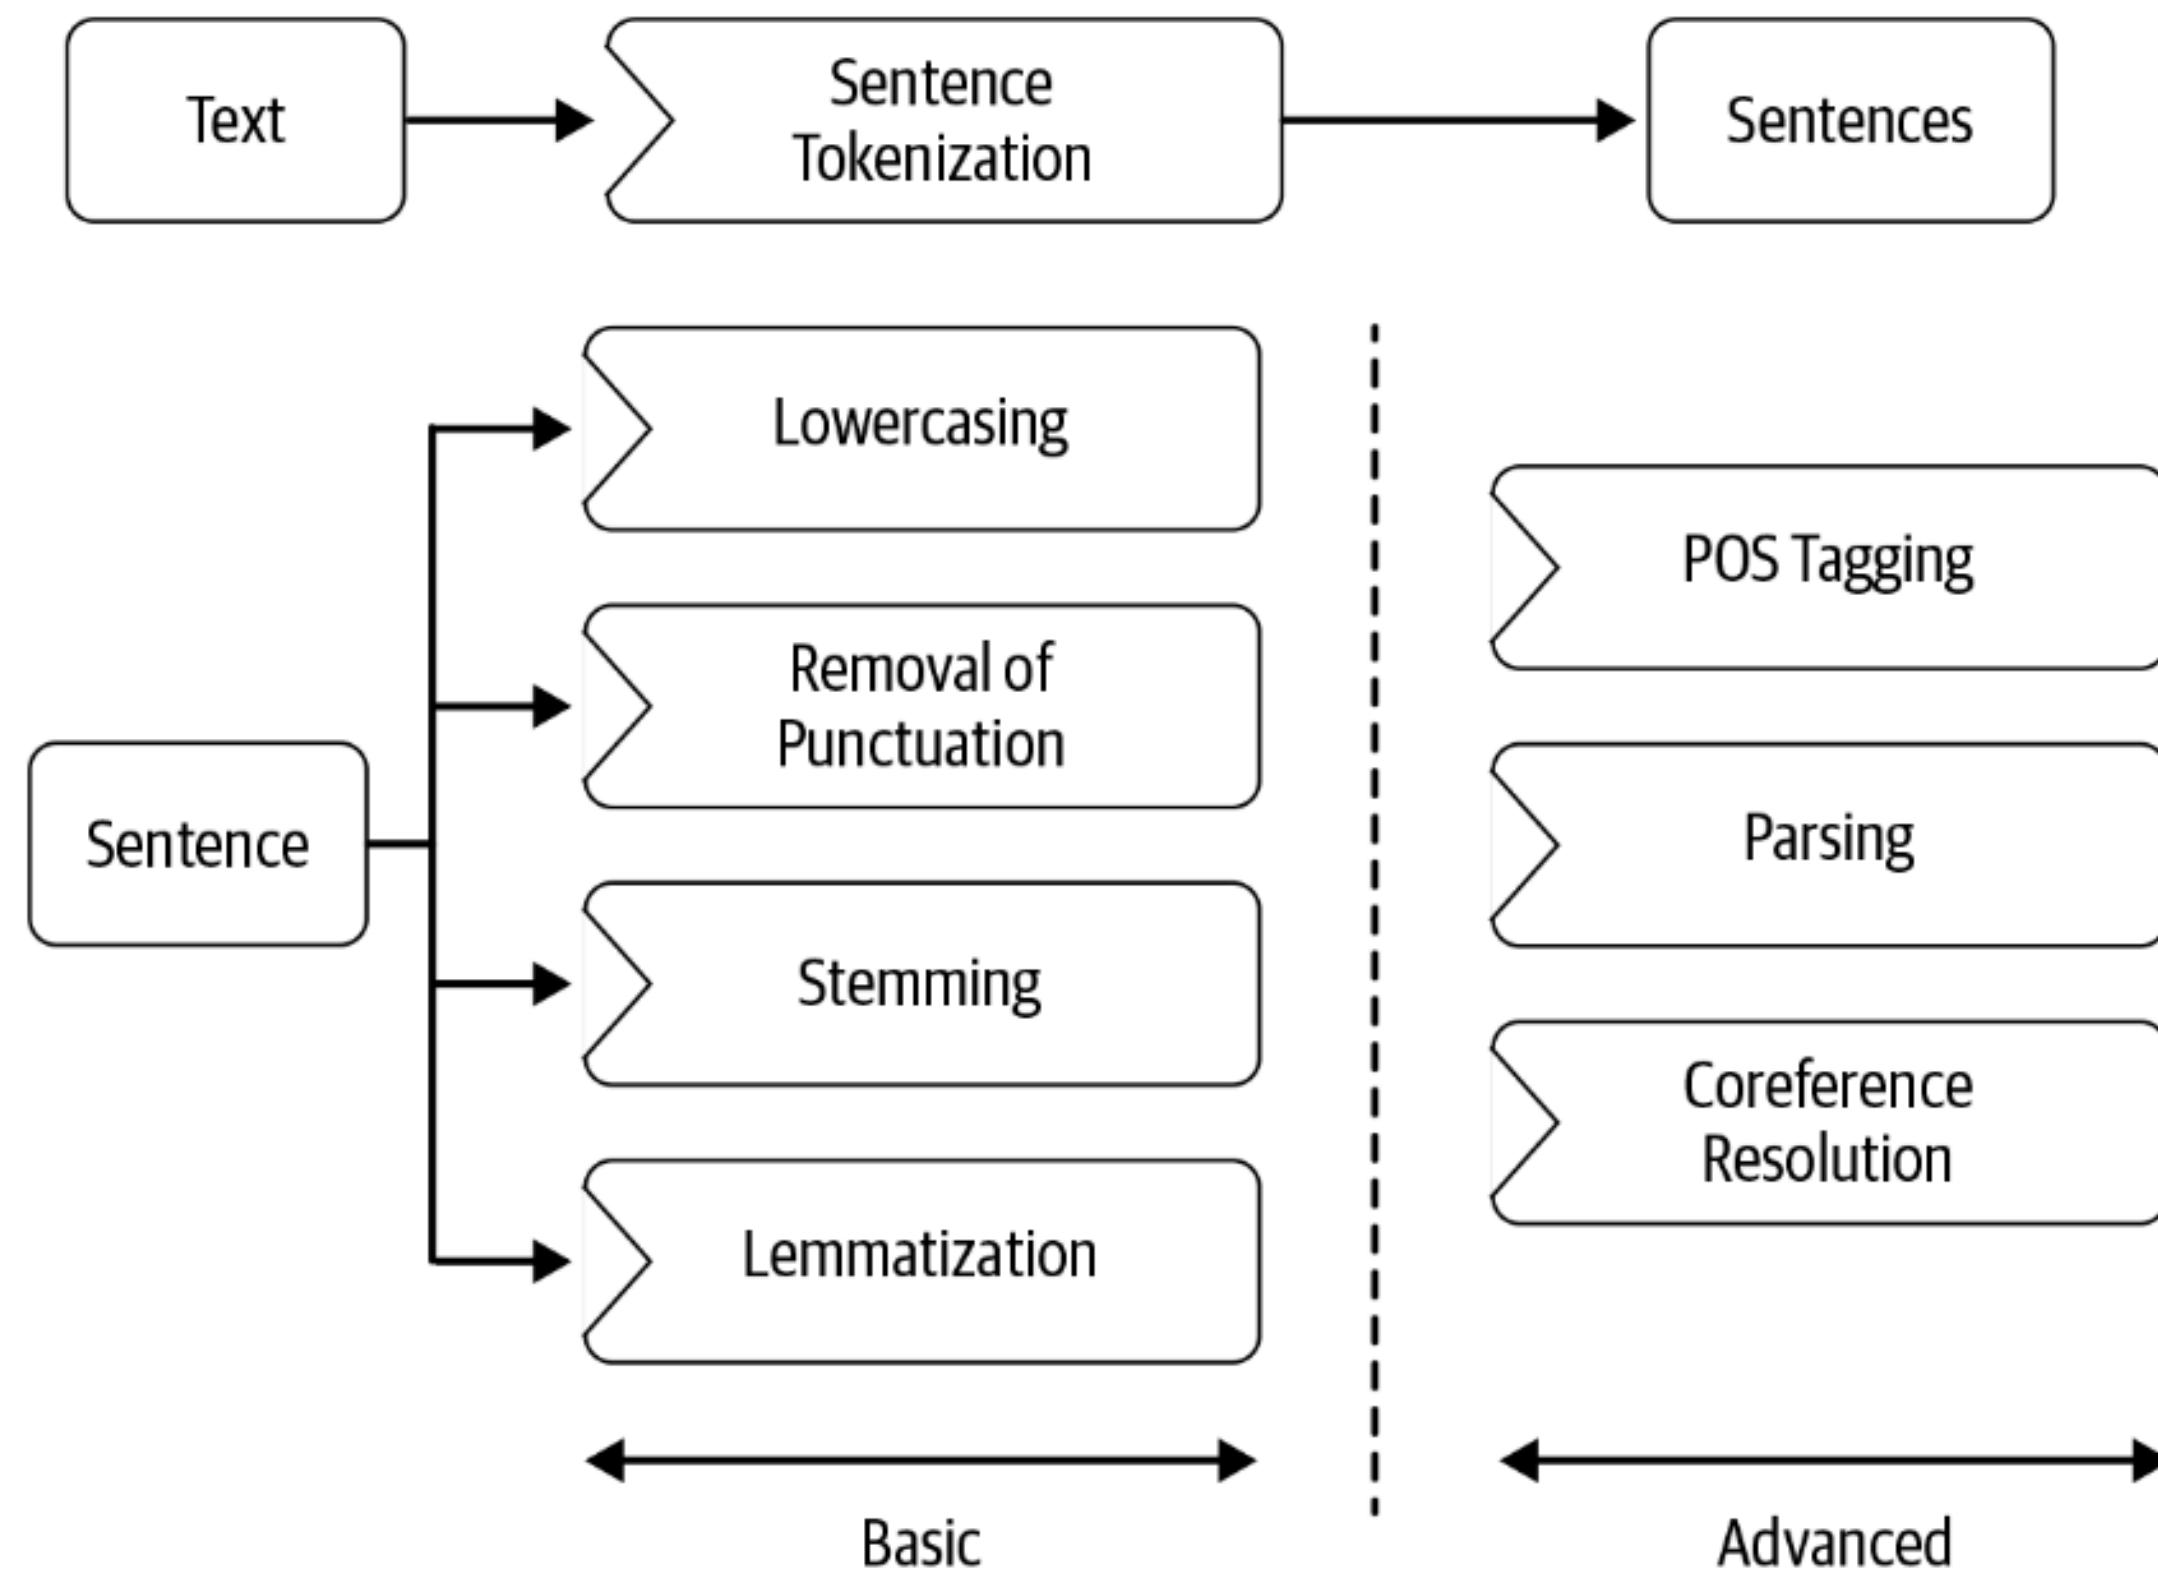
	$N$	$ \mathcal{V} $
<i>Switchboard</i>	2.4 million	20 thousand
<i>Shakespeare</i>	884 thousand	31 thousand
<i>Google N-gram</i>	1 trillion	13 million

# Pre-processing: summary



Summary of pre-processing steps. *From "Practical Language Processing, Figure 2-11, Chapter 2"*

# Pre-processing: summary



Summary of pre-processing steps. *From "Practical Language Processing, Figure 2-11, Chapter 2"*

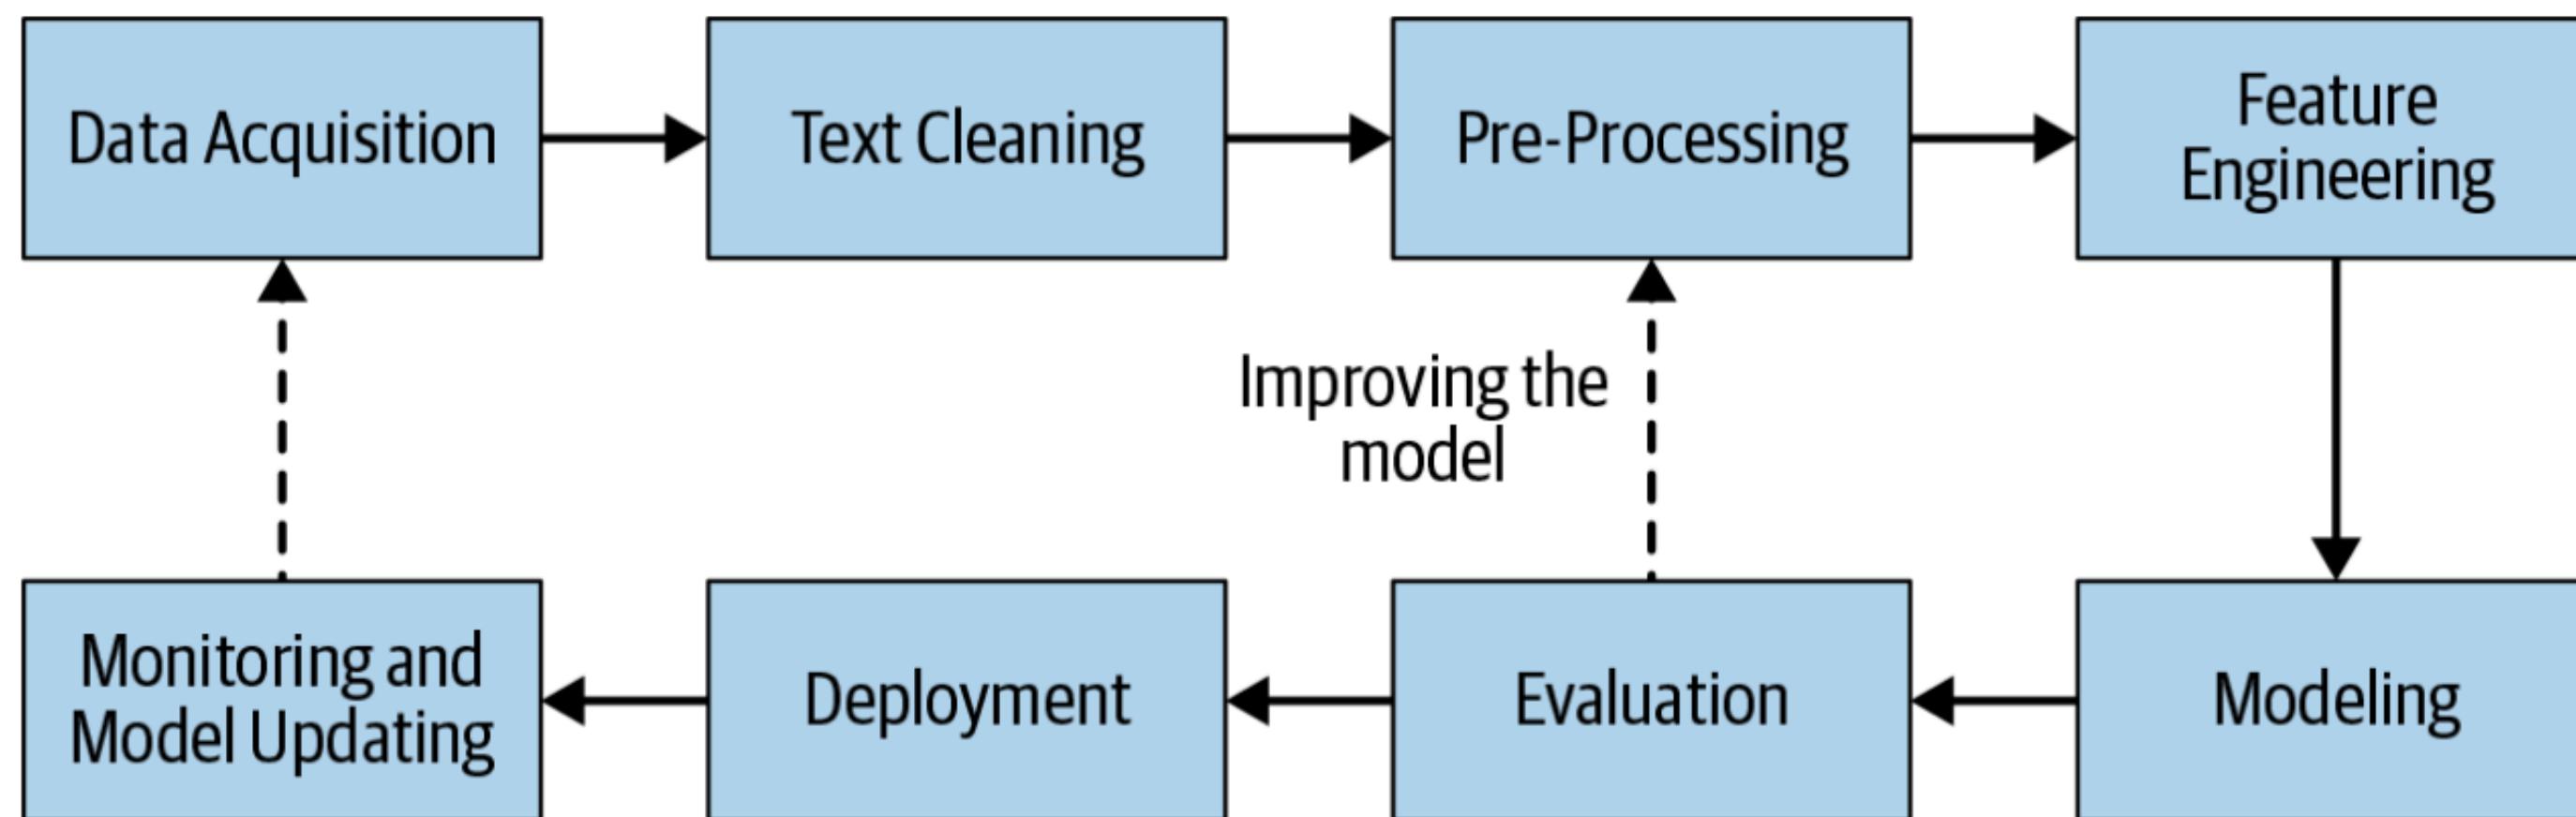
→ Possible advanced steps depending on the target task !

# A case study: classic methods for document representation and classification

---

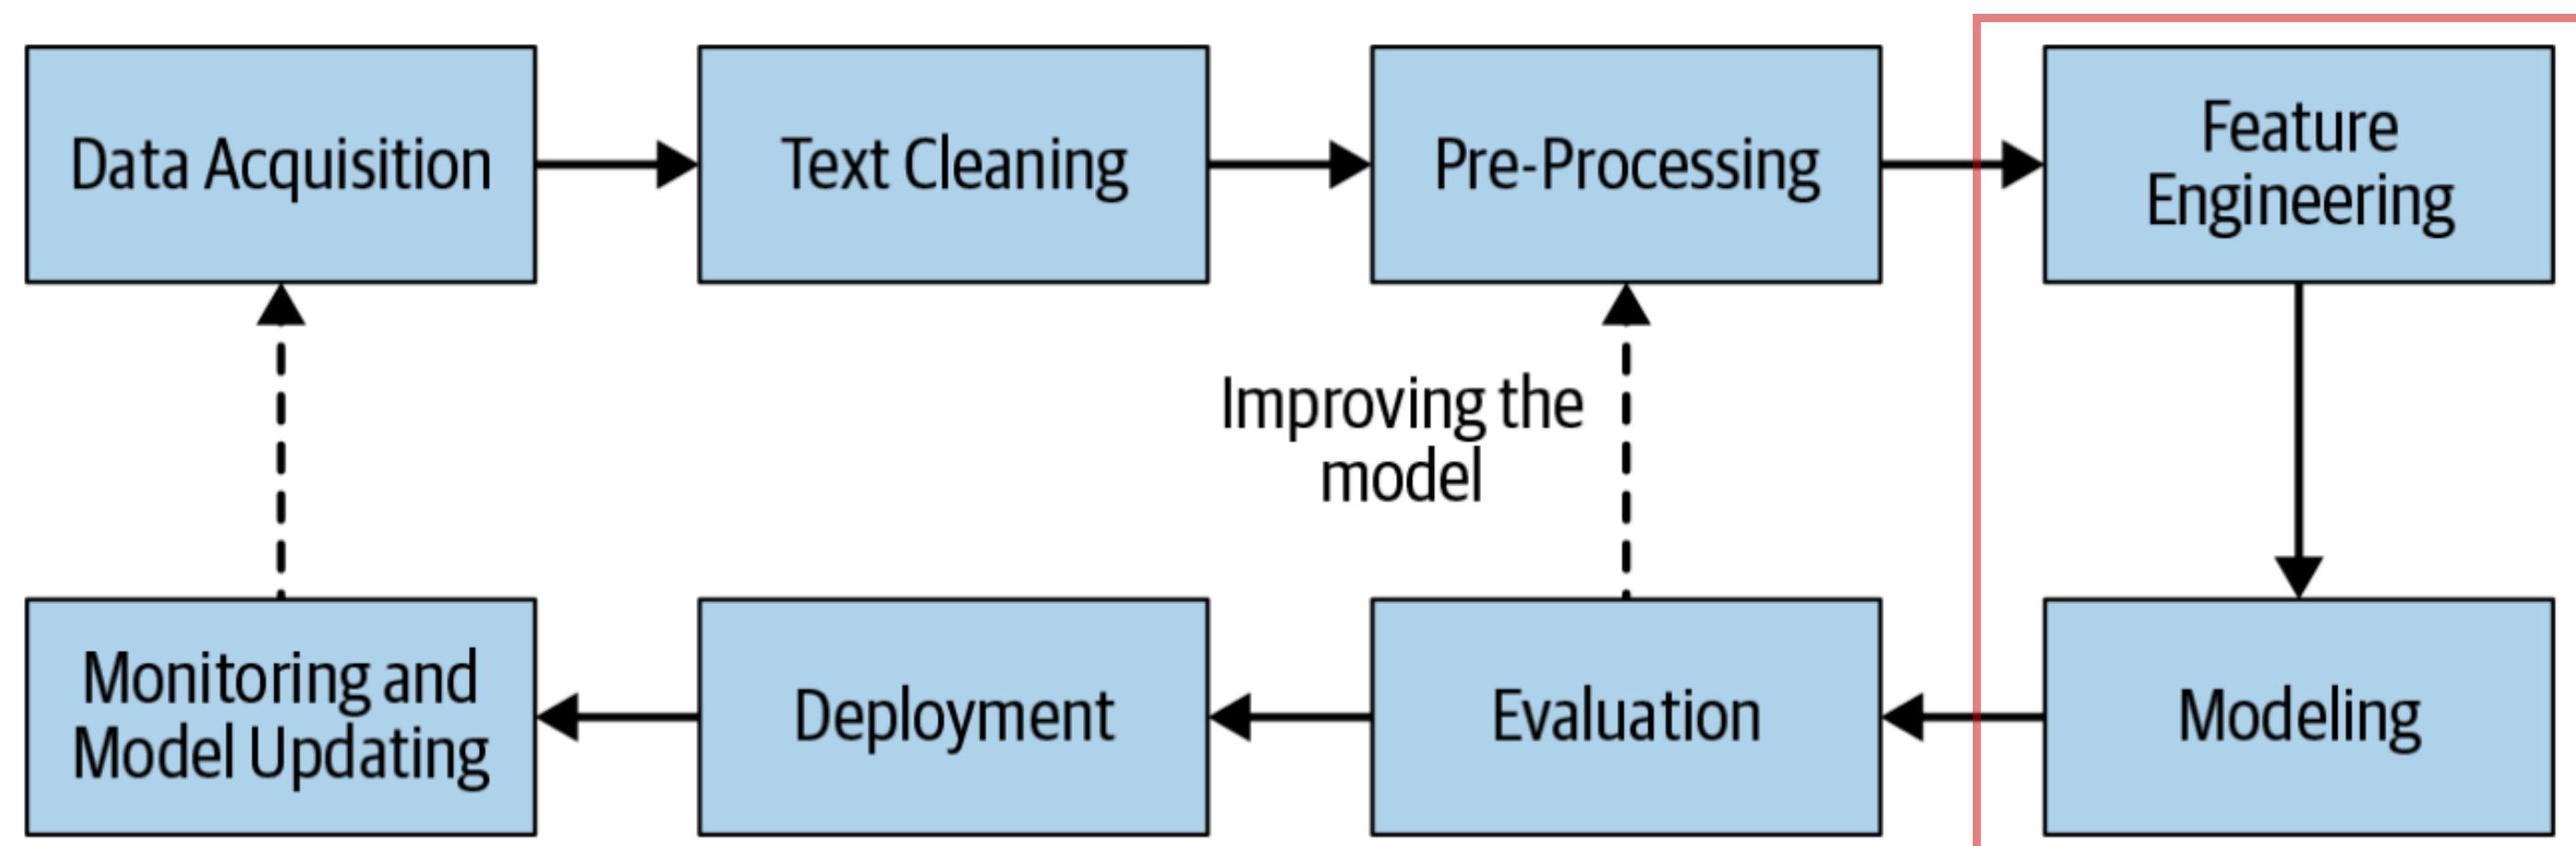
# A case study: classic methods for document representation and classification

---



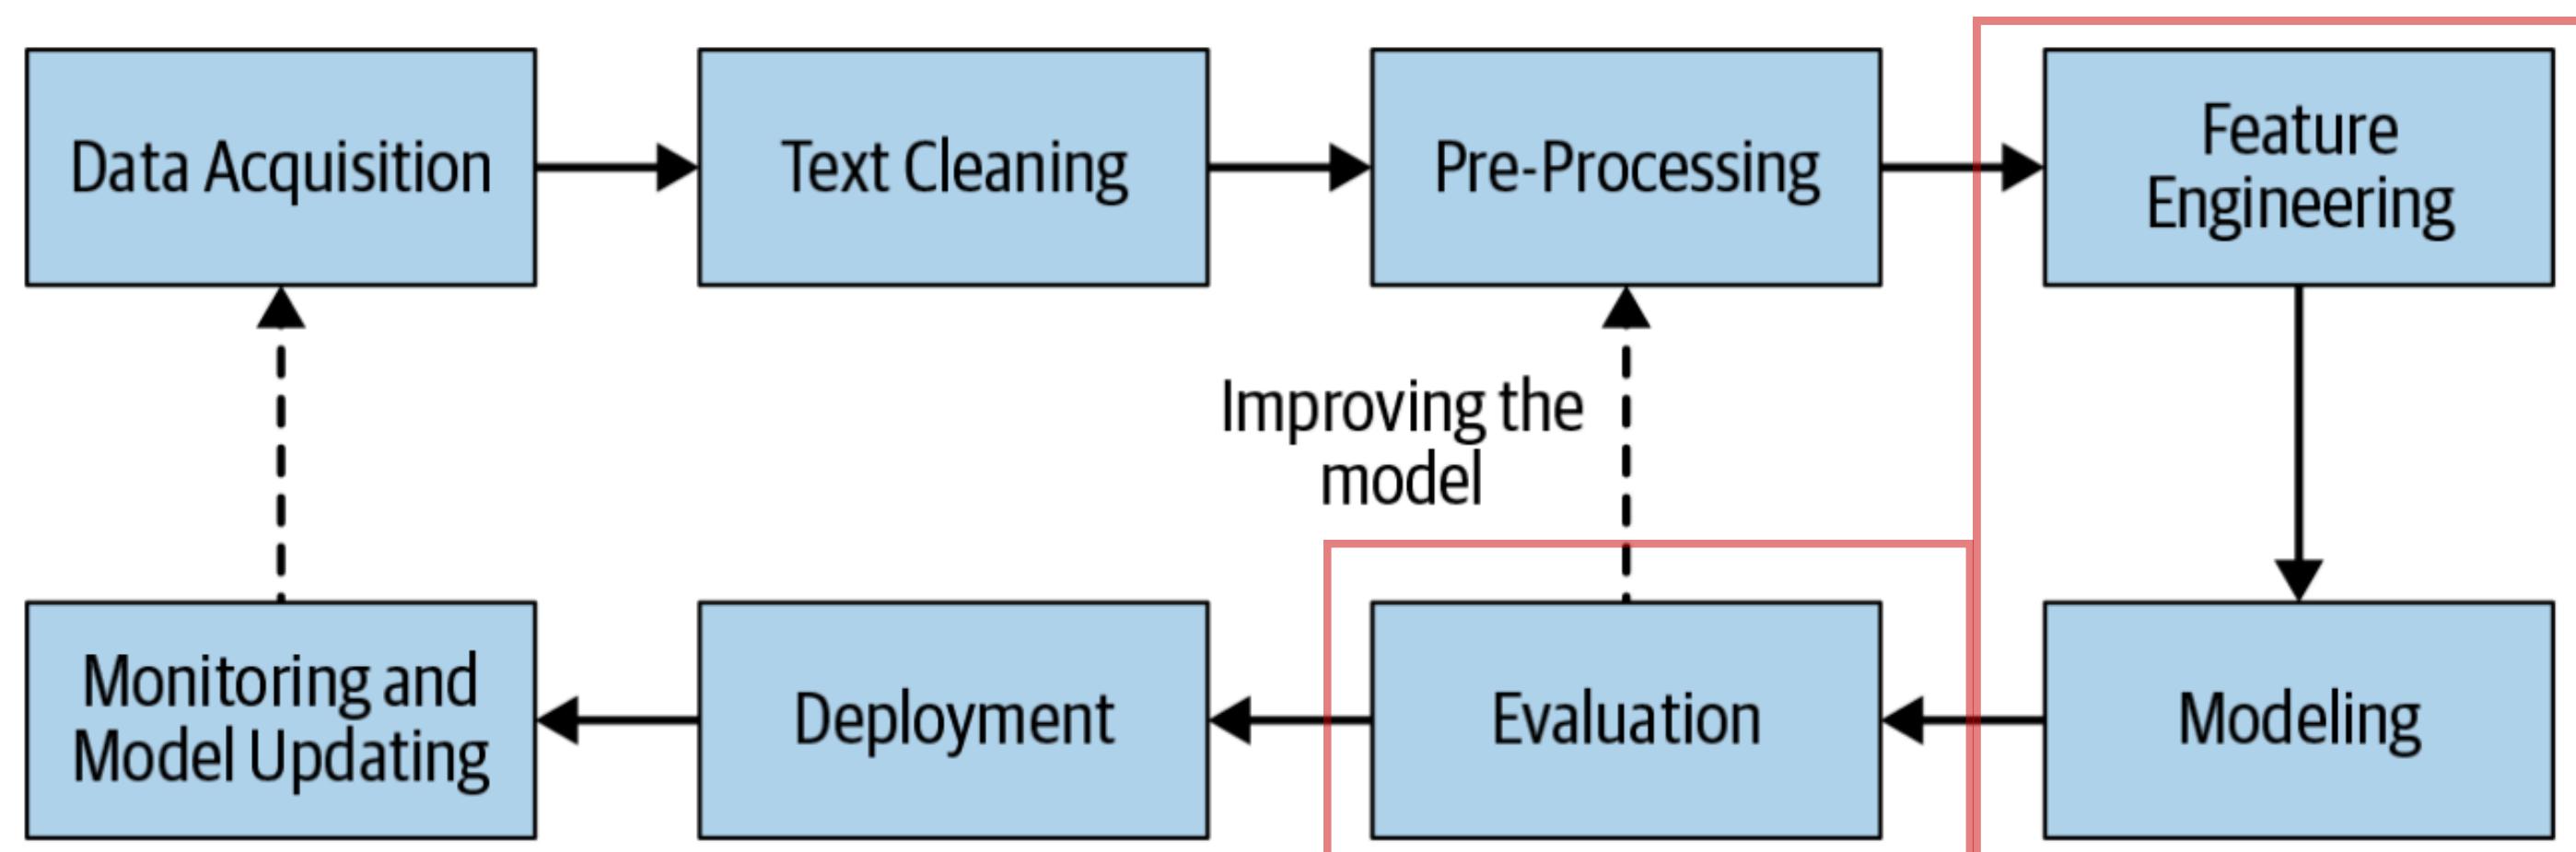
# A case study: classic methods for document representation and classification

---



# A case study: classic methods for document representation and classification

---



# Task: Text classification

---

Categorizing an instance (sentence, document..) into one or more **known** classes

# Task: Text classification

---

Categorizing an instance (sentence, document..) into one or more **known** classes

- A wide array of existing tasks, including *multi-class* and *multi-label*

# Task: Text classification

---

Categorizing an instance (sentence, document..) into one or more **known** classes

- A wide array of existing tasks, including *multi-class* and *multi-label*
- Can concern sentiment, topic, ... with very different granularities  
→ The label set can be structured (*hierarchical classification*)

# Task: Text classification

---

Categorizing an instance (sentence, document..) into one or more **known** classes

- A wide array of existing tasks, including *multi-class* and *multi-label*
- Can concern sentiment, topic, ... with very different granularities
  - The label set can be structured (*hierarchical classification*)
- Necessitate **labels**, hence **annotations**
  - Supervised (or semi-supervised) learning
  - Existing large scale datasets mined from user-generated content (online reviews), but it's not always the case
  - But: Human annotations are ambiguous
  - Label sets can be insufficiently defined

# Task: Text classification

---

Categorizing an instance (sentence, document..) into one or more **known** classes

- A wide array of existing tasks, including *multi-class* and *multi-label*
- Can concern sentiment, topic, ... with very different granularities
  - The label set can be structured (*hierarchical classification*)
- Necessitate **labels**, hence **annotations**
  - Supervised (or semi-supervised) learning
  - Existing large scale datasets mined from user-generated content (online reviews), but it's not always the case
  - But: Human annotations are ambiguous
  - Label sets can be insufficiently defined
- Well-defined and understood metrics, making evaluation straightforward

# Features: Document as bag-of-words

---

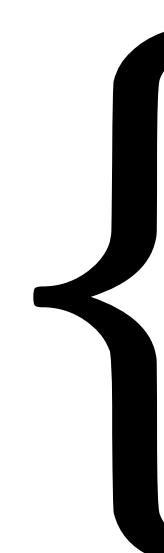
The simplest way to represent a document is as a **bag-of-word** (an *unordered set of words*, keeping only their frequency).

# Features: Document as bag-of-words

---

The simplest way to represent a document is as a **bag-of-word** (an *unordered set of words*, keeping only their frequency).

Assuming the following set of (short) documents: how to represent them ?

- 
- *I walked down the street*
  - *I walked down the avenue*
  - *I ran down the street*
  - *I walk down the city*
  - *I walk down the avenue*

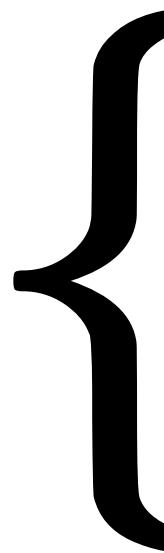
# Features: Document as bag-of-words

---

The simplest way to represent a document is as a **bag-of-word** (an *unordered set of words*, keeping only their frequency).

Assuming the following set of (short) documents: how to represent them ?

Compute the vocabulary  $\mathcal{V}$  and count !

- 
- *I walked down the street*
  - *I walked down the avenue*
  - *I ran down the street*
  - *I walk down the city*
  - *I walk down the avenue*

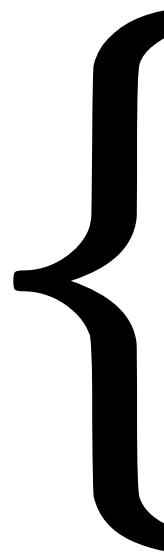
# Features: Document as bag-of-words

---

The simplest way to represent a document is as a **bag-of-word** (an *unordered set of words*, keeping only their frequency).

Assuming the following set of (short) documents: how to represent them ?

Compute the vocabulary  $\mathcal{V}$  and count !

- 
- *I walked down the street*
  - *I walked down the avenue*
  - *I ran down the street*
  - *I walk down the city*
  - *I walk down the avenue*

$$\rightarrow \mathcal{V} = \{\text{I, the, down, walked, street, avenue, walk, ran, city}\}$$

# Features: Document as bag-of-words

---

The simplest way to represent a document is as a **bag-of-word** (an *unordered set of words*, keeping only their frequency).

Assuming the following set of (short) documents: how to represent them ?

Compute the vocabulary  $\mathcal{V}$  and count !

- *I walked down the street*
- *I walked down the avenue*
- *I ran down the street*
- *I walk down the city*
- *I walk down the avenue*

$$\rightarrow \mathcal{V} = \{\text{I, the, down, walked, street, avenue, walk, ran, city}\}$$

$\rightarrow \mathbf{T} =$   
*(Term-document matrix)*

	I	the	down	walked	street	avenue	walk	ran	city
Doc_1	1	1	1	1	1	0	0	0	0
Doc_2	1	1	1	1	0	1	0	0	0
Doc_3	1	1	1	0	1	0	0	1	0
Doc_4	1	1	1	0	0	0	1	0	1
Doc_5	1	1	1	0	0	1	1	0	0

# Document representation: motivation

---

The **Bag-of-word** is a document model counting words

- Assumes that position does not matter ... hence indifferent to syntax and semantics

# Document representation: motivation

---

The **Bag-of-word** is a document model counting words

- Assumes that position does not matter ... hence indifferent to syntax and semantics
- Main goal: text classification (sentiment, spam ...) !

# Document representation: motivation

---

The **Bag-of-word** is a document model counting words

- Assumes that position does not matter ... hence indifferent to syntax and semantics
- Main goal: text classification (sentiment, spam ...) !
  - We can use *rules* based on words ...  
→ with *Sentiwordnet*: Each Synset is associated to a positive and negative score.

*SENTIWORDNET 3.0: An Enhanced Lexical Resource for Sentiment Analysis and Opinion Mining*  
(Baccianella et al., 2010)

# Document representation: motivation

---

The **Bag-of-word** is a document model counting words

- Assumes that position does not matter ... hence indifferent to syntax and semantics
- Main goal: text classification (sentiment, spam ...) !
  - We can use *rules* based on words ...  
→ with *Sentiwordnet*: Each Synset is associated to a positive and negative score.

*SENTIWORDNET 3.0: An Enhanced Lexical Resource for Sentiment Analysis and Opinion Mining*  
(Baccianella et al., 2010)

- ...or learn a classifier model that will use *word frequencies* as features → *Naïve Bayes* is the simplest, assuming independance between words

# Document representation: motivation

---

The **Bag-of-word** is a document model counting words

- Assumes that position does not matter ... hence indifferent to syntax and semantics
- Main goal: text classification (sentiment, spam ...) !
  - We can use *rules* based on words ...  
→ with *Sentiwordnet*: Each Synset is associated to a positive and negative score.

*SENTIWORDNET 3.0: An Enhanced Lexical Resource for Sentiment Analysis and Opinion Mining*  
(Baccianella et al., 2010)

- ...or learn a classifier model that will use *word frequencies* as features → *Naïve Bayes* is the simplest, assuming independance between words
- Also useful for document clustering, information retrieval.... why ?

# Model: Naïve Bayes

---

**Multinomial naïve Bayes classifier:** a *generative* (why ?) linear classifier that naïvely assumes that features are independant

- Goal: for a document  $d$  return the class  $\hat{c}$  with maximum a posteriori probability among classes:  $\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} P(c|d)$

# Model: Naïve Bayes

---

**Multinomial naive Bayes classifier:** a *generative* (why ?) linear classifier that naïvely assumes that features are independant

- Goal: for a document  $d$  return the class  $\hat{c}$  with maximum a posteriori probability among classes:  $\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} P(c|d)$
- Applying **Baye's rule** and the **independance assumption**, and noting  $d = (w_1, \dots, w_n)$  we get a (*prior*  $\times$  *likelihood*) decomposition:

$$\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} \left[ \mathbb{P}(c) \prod_{i=1}^n \mathbb{P}(w_i|c) \right]$$

- *Derivation ?*

# Model: Naïve Bayes

---

**Multinomial naive Bayes classifier:** a *generative* (why ?) linear classifier that naïvely assumes that features are independant

- Goal: for a document  $d$  return the class  $\hat{c}$  with maximum a posteriori probability among classes:  $\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} P(c|d)$
- Applying **Baye's rule** and the **independance assumption**, and noting  $d = (w_1, \dots, w_n)$  we get a (*prior*  $\times$  *likelihood*) decomposition:

$$\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} \left[ \mathbb{P}(c) \prod_{i=1}^n \mathbb{P}(w_i|c) \right]$$

- *Derivation ?*
- Idea: we compute  $\mathbb{P}(w|c)$  as the frequency of  $w$  among all documents  $d \in \mathcal{D}$  of class  $c$ :

$$\mathbb{P}(w|c) = \frac{\operatorname{count}_{\mathcal{D}}(w, c)}{\sum_{w' \in \mathcal{V}} \operatorname{count}_{\mathcal{D}}(w', c)}$$

# Model: Naïve Bayes

---

**Multinomial naive Bayes classifier:** a *generative* (why ?) linear classifier that naïvely assumes that features are independant

- Goal: for a document  $d$  return the class  $\hat{c}$  with maximum a posteriori probability among classes:  $\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} P(c|d)$
- Applying **Baye's rule** and the **independance assumption**, and noting  $d = (w_1, \dots, w_n)$  we get a (*prior*  $\times$  *likelihood*) decomposition:

$$\hat{c} = \operatorname{argmax}_{c \in \mathcal{C}} \left[ \mathbb{P}(c) \prod_{i=1}^n \mathbb{P}(w_i|c) \right]$$

- *Derivation* ?
- Idea: we compute  $\mathbb{P}(w|c)$  as the frequency of  $w$  among all documents  $d \in \mathcal{D}$  of class  $c$ :

$$\mathbb{P}(w|c) = \frac{\operatorname{count}_{\mathcal{D}}(w, c)}{\sum_{w' \in \mathcal{V}} \operatorname{count}_{\mathcal{D}}(w', c)}$$

- This is Maximum Likelihood Estimation (MLE): *Demonstration* ?

# Classification with Naïve Bayes

---

- It is legitimate (and practical !) to represent a document  $d$  with its bag-of-word representation  $\mathbf{d}$ !
  - Practical points: use log-probabilities (why ?), use *smoothing*

# Classification with Naïve Bayes

---

- It is legitimate (and practical !) to represent a document  $d$  with its bag-of-word representation  $\mathbf{d}$ !
  - Practical points: use log-probabilities (why ?), use *smoothing*
- **Training Algorithm:** Given data  $\mathcal{D}$  and classes  $\mathcal{C}$ 
  - Create the vocabulary  $\mathcal{V}$  from documents  $d \in \mathcal{D}$
  - From  $\mathcal{D}$ : For each class  $c \in \mathcal{C}$ : compute the log-prior  $\log \mathbb{P}(c)$ 
    - For each word  $w \in \mathcal{V}$  compute the log-likelihood  $\log \mathbb{P}(w|c)$

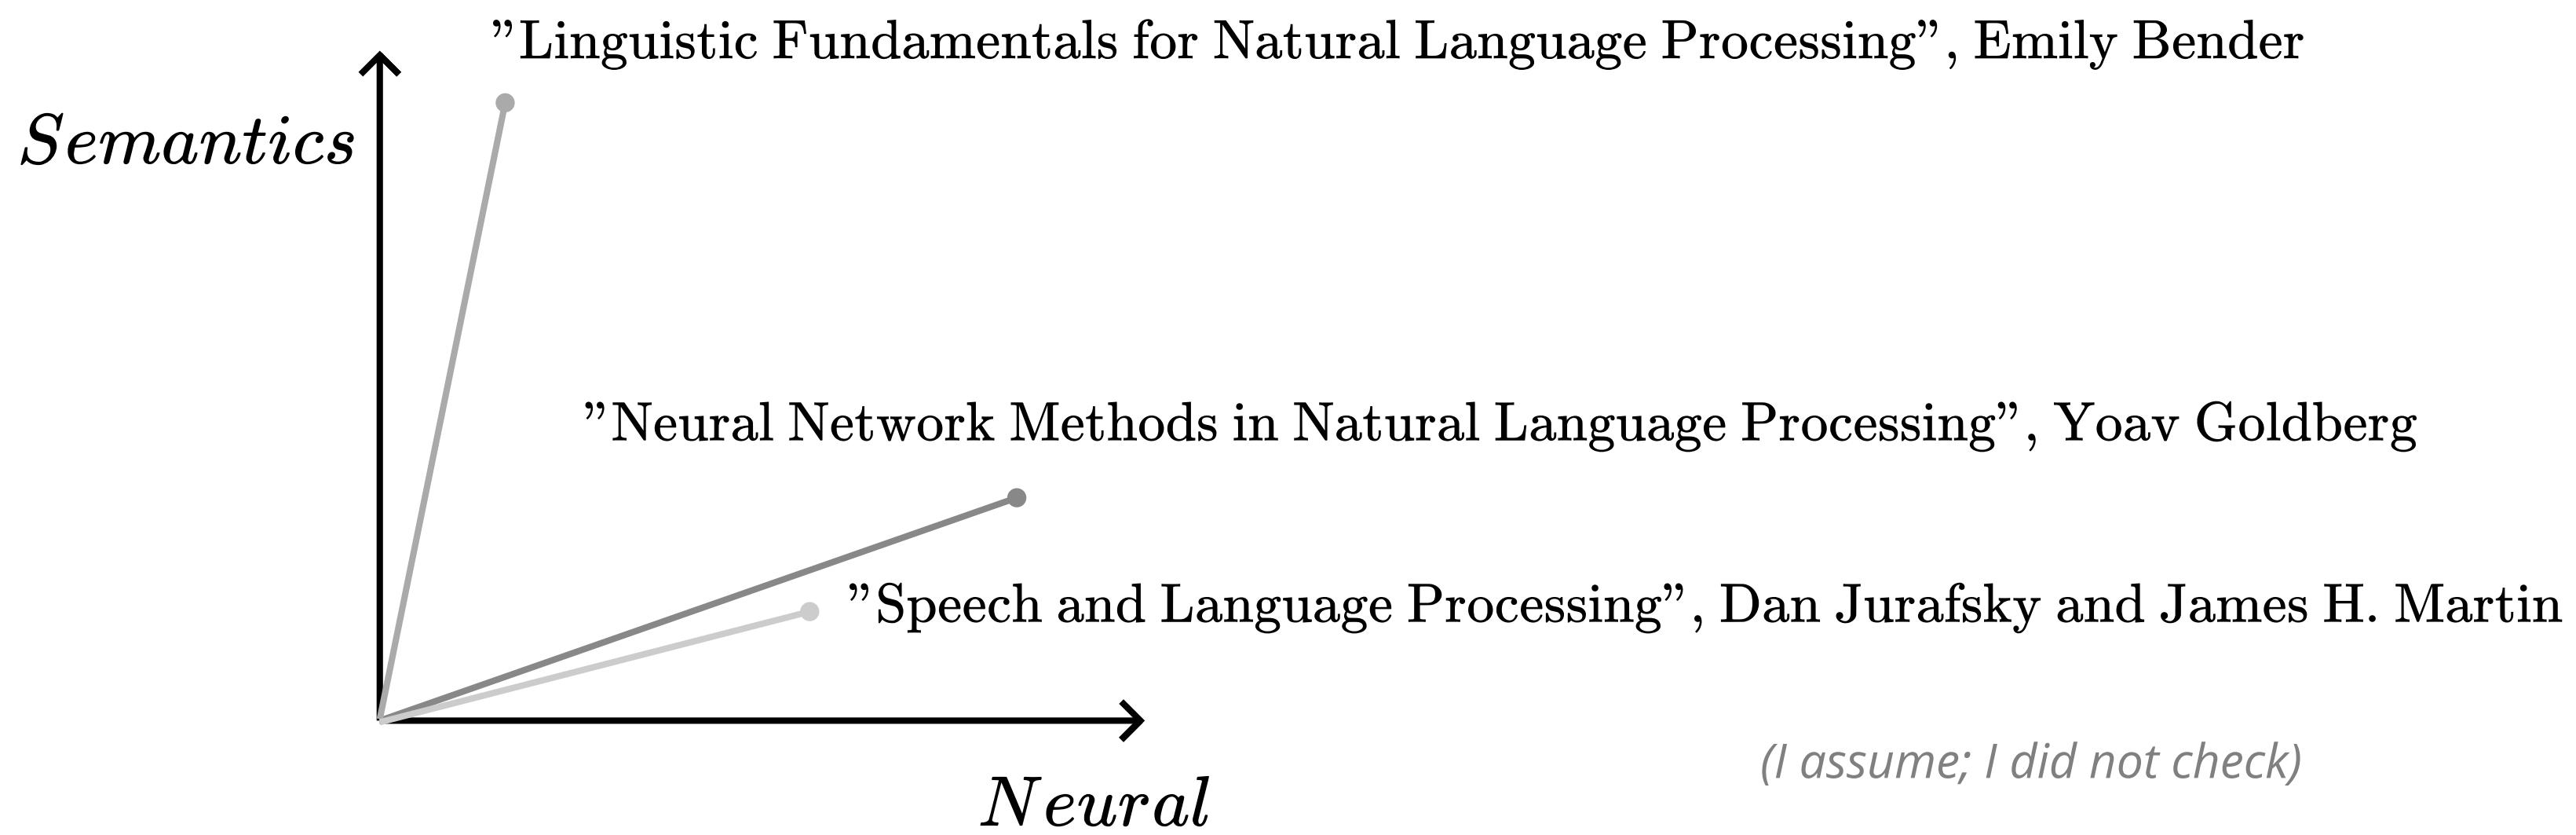
# Classification with Naïve Bayes

---

- It is legitimate (and practical !) to represent a document  $d$  with its bag-of-word representation  $\mathbf{d}$ !
  - Practical points: use log-probabilities (why ?), use *smoothing*
- **Training Algorithm:** Given data  $\mathcal{D}$  and classes  $\mathcal{C}$ 
  - Create the vocabulary  $\mathcal{V}$  from documents  $d \in \mathcal{D}$
  - From  $\mathcal{D}$ : For each class  $c \in \mathcal{C}$ : compute the log-prior  $\log \mathbb{P}(c)$ 
    - For each word  $w \in \mathcal{V}$  compute the log-likelihood  $\log \mathbb{P}(w|c)$
- **Inference Algorithm:** Given document  $d$  to classify:
  - For each class  $c \in \mathcal{C}$ :  $S(c) = \log \mathbb{P}(c)$ 
    - For each position  $i \in d$ :
      - If  $w_i \in \mathcal{V}$ :  $S(c) = S(c) + \log \mathbb{P}(w_i|c)$
  - Return  $\text{argmax}_{c \in \mathcal{C}} S(c)$

# Document as Vectors

---



- Words are **dimension** of documents vectors
- You can visualize vectors in a particular set of dimensions of your choosing
- Vectors should be *similar* for documents that are related

But what does "similar" mean here ?

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

- Based on the dot product - but alone, it favors long documents, with more words and higher values.

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

- Based on the dot product - but alone, it favors long documents, with more words and higher values.
- Normalizing the dot product gives us the cosine of the angle between the two vectors:

$$\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = \frac{\mathbf{d}_1 \cdot \mathbf{d}_2}{\|\mathbf{d}_1\| \|\mathbf{d}_2\|}$$

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

- Based on the dot product - but alone, it favors long documents, with more words and higher values.
- Normalizing the dot product gives us the cosine of the angle between the two vectors:

$$\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = \frac{\mathbf{d}_1 \cdot \mathbf{d}_2}{\|\mathbf{d}_1\| \|\mathbf{d}_2\|}$$

- Values range in  $[-1, 1]$ ; however, with frequencies as features, the similarity is always positive

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

- Based on the dot product - but alone, it favors long documents, with more words and higher values.
- Normalizing the dot product gives us the cosine of the angle between the two vectors:

$$\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = \frac{\mathbf{d}_1 \cdot \mathbf{d}_2}{\|\mathbf{d}_1\| \|\mathbf{d}_2\|}$$

- Values range in  $[-1, 1]$ ; however, with frequencies as features, the similarity is always positive
- $\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = 0 \rightarrow$  the documents have no words in common

# Similarity between documents: cosine

---

The **cosine** of the angle between the document vectors is the most common similarity metric in NLP

- Based on the dot product - but alone, it favors long documents, with more words and higher values.
- Normalizing the dot product gives us the cosine of the angle between the two vectors:

$$\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = \frac{\mathbf{d}_1 \cdot \mathbf{d}_2}{\|\mathbf{d}_1\| \|\mathbf{d}_2\|}$$

- Values range in  $[-1, 1]$ ; however, with frequencies as features, the similarity is always positive
- $\text{cosine}(\mathbf{d}_1, \mathbf{d}_2) = 0 \rightarrow$  the documents have no words in common

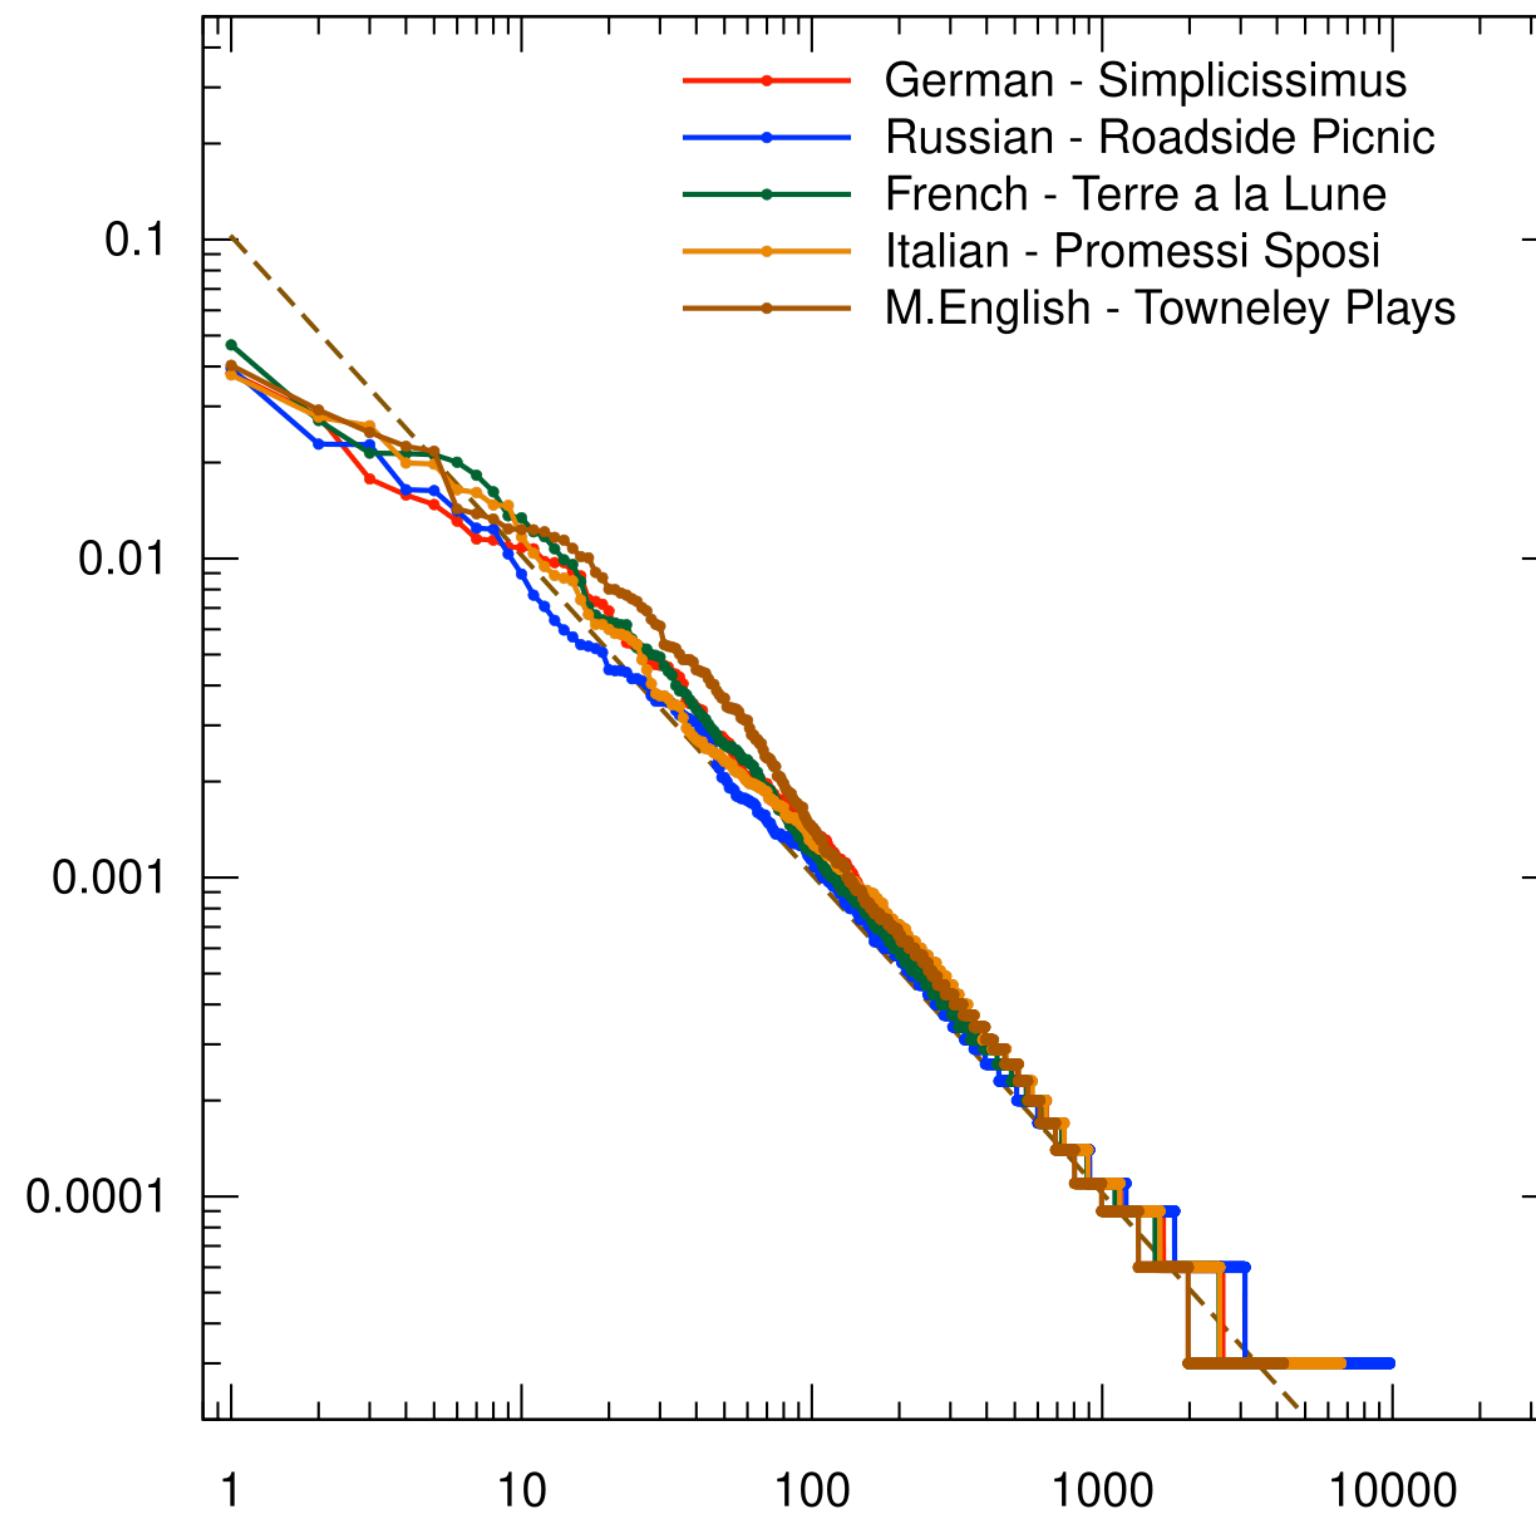
Still, frequency is not the best measure of association between words:

- It is **skewed**  $\rightarrow$  Zipf's law
- Very frequent words are rarely the most useful for classification

# Difficulty 1: Zipf's law

---

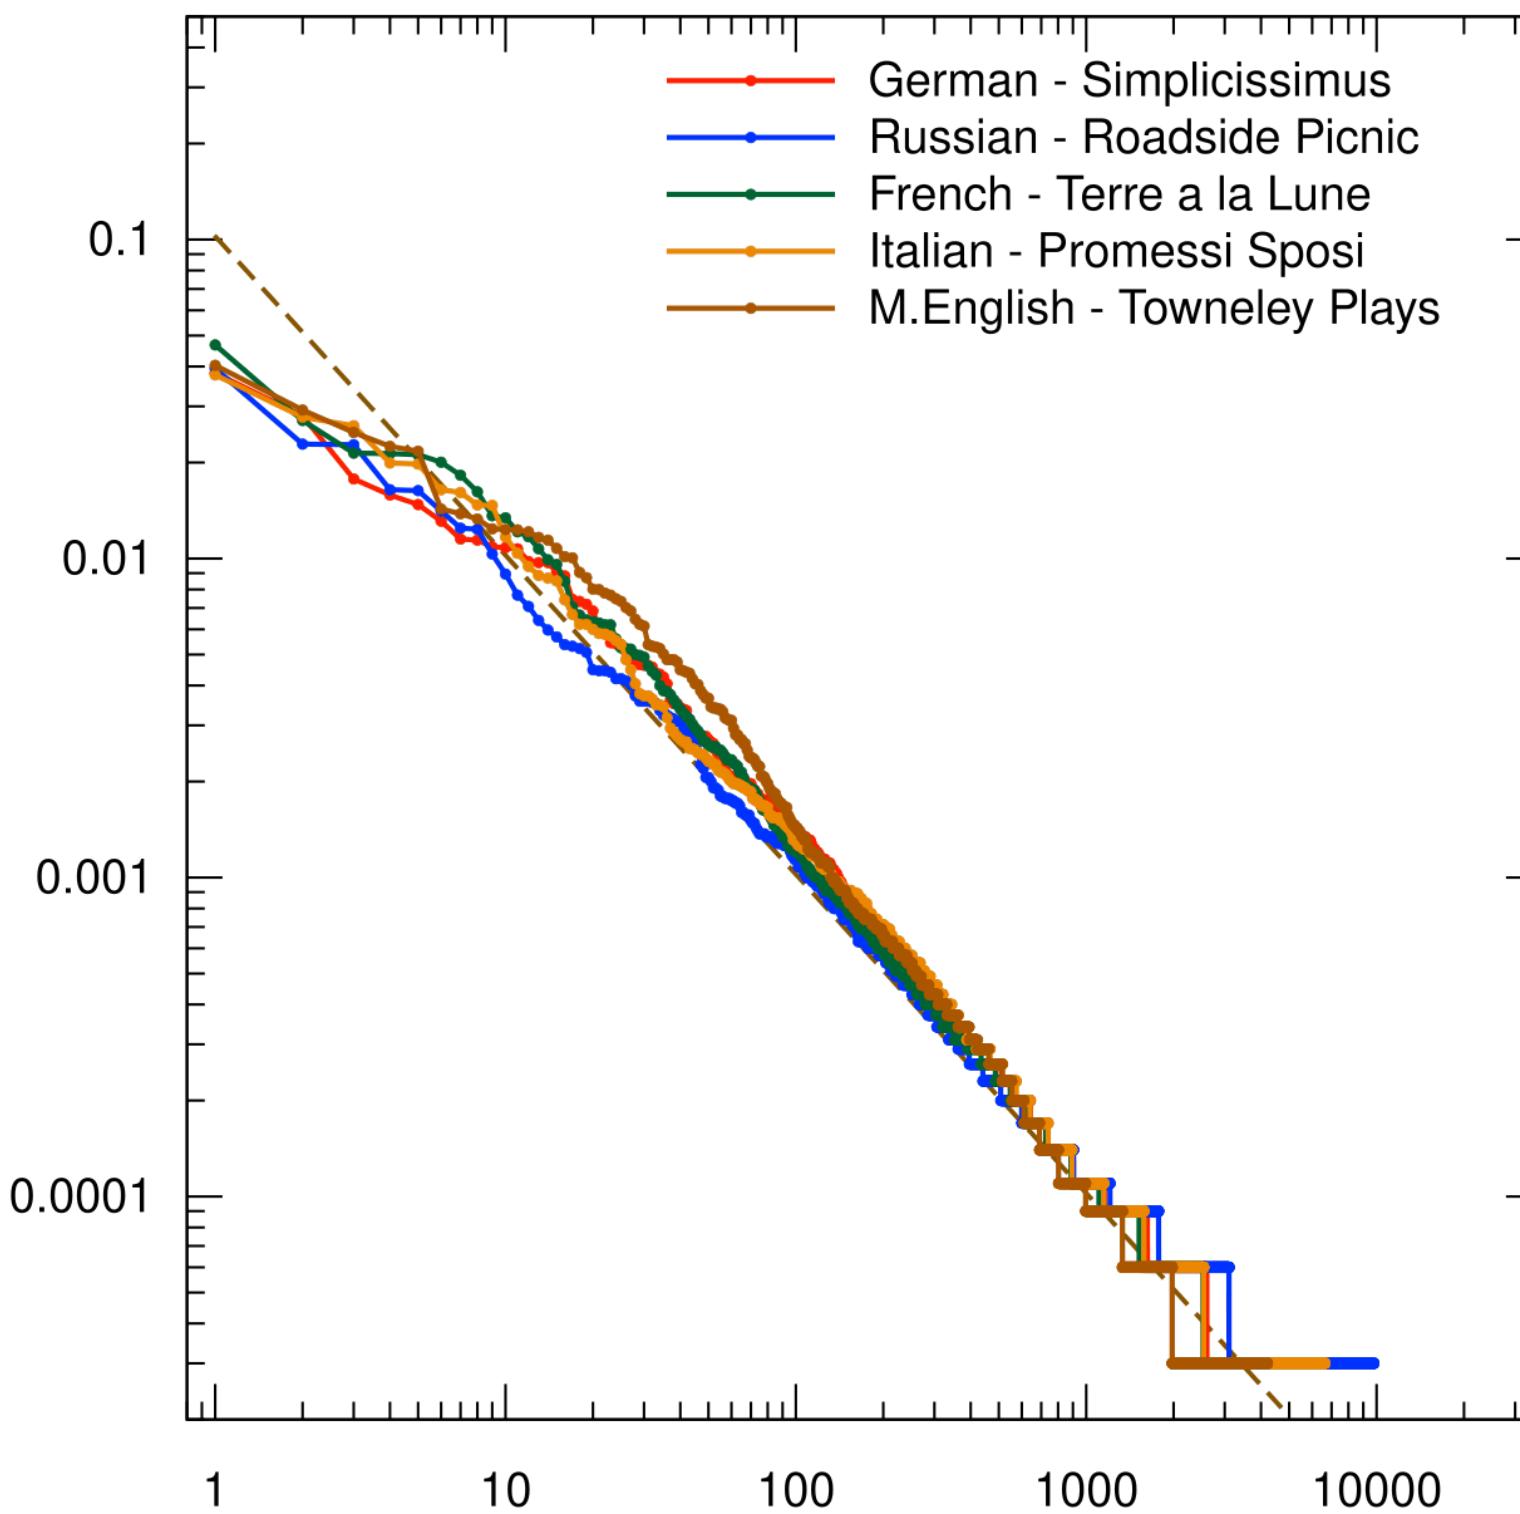
**Sparsity:** tied to **Zipf's Law:**  $frequency \propto 1/rank$



# Difficulty 1: Zipf's law

---

**Sparsity:** tied to **Zipf's Law:**  $frequency \propto 1/rank$



- Zipf's law seems to hold for most natural languages and many language-related phenomena
  - Examples: *meaning-frequency law, law of abbreviation*

# Features: TF-IDF Representations

---

What would be a better way than directly removing frequent-but-not-significant words (called *stopwords*) ?

# Features: TF-IDF Representations

---

What would be a better way than directly removing frequent-but-not-significant words (called *stopwords*) ?

- Idea 1: instead of using count  $c(w, d)$  of word  $w$  in document  $d$ , define *term frequency* as  $\text{TF}(w, d) = \log_{10}(c(w, d) + 1)$

# Features: TF-IDF Representations

---

What would be a better way than directly removing frequent-but-not-significant words (called *stopwords*) ?

- Idea 1: instead of using count  $c(w, d)$  of word  $w$  in document  $d$ , define *term frequency* as  $\text{TF}(w, d) = \log_{10}(c(w, d) + 1)$
- Idea 2: give higher weight to words that occur in only a few documents, using their *inverse document frequency*. Noting  $cd(w)$  the count of documents  $w$  appears in and  $N$  the total number of documents,

$$\text{IDF}(w) = \log_{10} \left( \frac{N}{cd(w)} \right)$$

# Features: TF-IDF Representations

---

What would be a better way than directly removing frequent-but-not-significant words (called *stopwords*) ?

- Idea 1: instead of using count  $c(w, d)$  of word  $w$  in document  $d$ , define *term frequency* as  $\text{TF}(w, d) = \log_{10}(c(w, d) + 1)$
- Idea 2: give higher weight to words that occur in only a few documents, using their *inverse document frequency*. Noting  $cd(w)$  the count of documents  $w$  appears in and  $N$  the total number of documents,

$$\text{IDF}(w) = \log_{10} \left( \frac{N}{cd(w)} \right)$$

The weight given to word  $w$  in document  $d$  is  $\text{TF}(w, d) \times \text{IDF}(w)$ :

- What happens if a word is present in every document ?

# Model: Logistic regression

---

A *discriminative* linear classifier: learns directly  $\mathbb{P}(c|d)$  through computing a **linear score** and applying a **logistic function**.

- Binary case: for a set of documents  $d \in \mathcal{D}$  represented by vectors  $\mathbf{d}$  learn a vector  $\mathbf{w}$  and a bias  $b$  maximizing the likelihood of making a good classification into  $c = 1$  or  $c = 0$ .

# Model: Logistic regression

---

A *discriminative* linear classifier: learns directly  $\mathbb{P}(c|d)$  through computing a **linear score** and applying a **logistic function**.

- Binary case: for a set of documents  $d \in \mathcal{D}$  represented by vectors  $\mathbf{d}$  learn a vector  $\mathbf{w}$  and a bias  $b$  maximizing the likelihood of making a good classification into  $c = 1$  or  $c = 0$ .
- The probability  $\mathbb{P}(c = 1)$  is obtained by applying the sigmoid to the *scalar product plus intercept*:

$$\mathbb{P}(c = 1) = \sigma(\mathbf{w} \cdot \mathbf{d} + b)$$

# Model: Logistic regression

---

A *discriminative* linear classifier: learns directly  $\mathbb{P}(c|d)$  through computing a **linear score** and applying a **logistic function**.

- Binary case: for a set of documents  $d \in \mathcal{D}$  represented by vectors  $\mathbf{d}$  learn a vector  $\mathbf{w}$  and a bias  $b$  maximizing the likelihood of making a good classification into  $c = 1$  or  $c = 0$ .
- The probability  $\mathbb{P}(c = 1)$  is obtained by applying the sigmoid to the *scalar product plus intercept*:

$$\mathbb{P}(c = 1) = \sigma(\mathbf{w} \cdot \mathbf{d} + b)$$

- We want to maximize the likelihood of our data by minimizing the **cross-entropy** between true and predicted classes  $\hat{c}$  and  $c$ :

$$L(\hat{c}, c) = -\log \mathbb{P}(c|d) = -[c \log \hat{c} + (1 - c) \log(1 - \hat{c})]$$

# Model: Logistic regression

---

A *discriminative* linear classifier: learns directly  $\mathbb{P}(c|d)$  through computing a **linear score** and applying a **logistic function**.

- Binary case: for a set of documents  $d \in \mathcal{D}$  represented by vectors  $\mathbf{d}$  learn a vector  $\mathbf{w}$  and a bias  $b$  maximizing the likelihood of making a good classification into  $c = 1$  or  $c = 0$ .

- The probability  $\mathbb{P}(c = 1)$  is obtained by applying the sigmoid to the *scalar product plus intercept*:

$$\mathbb{P}(c = 1) = \sigma(\mathbf{w} \cdot \mathbf{d} + b)$$

- We want to maximize the likelihood of our data by minimizing the **cross-entropy** between true and predicted classes  $\hat{c}$  and  $c$ :

$$L(\hat{c}, c) = -\log \mathbb{P}(c|d) = -[c \log \hat{c} + (1 - c) \log(1 - \hat{c})]$$

- Here, the training is made through **gradient descent**: we minimize that loss function by finding iteratively the direction in which the loss decreases the most and updating the weights accordingly

# Classification with logistic regression

---

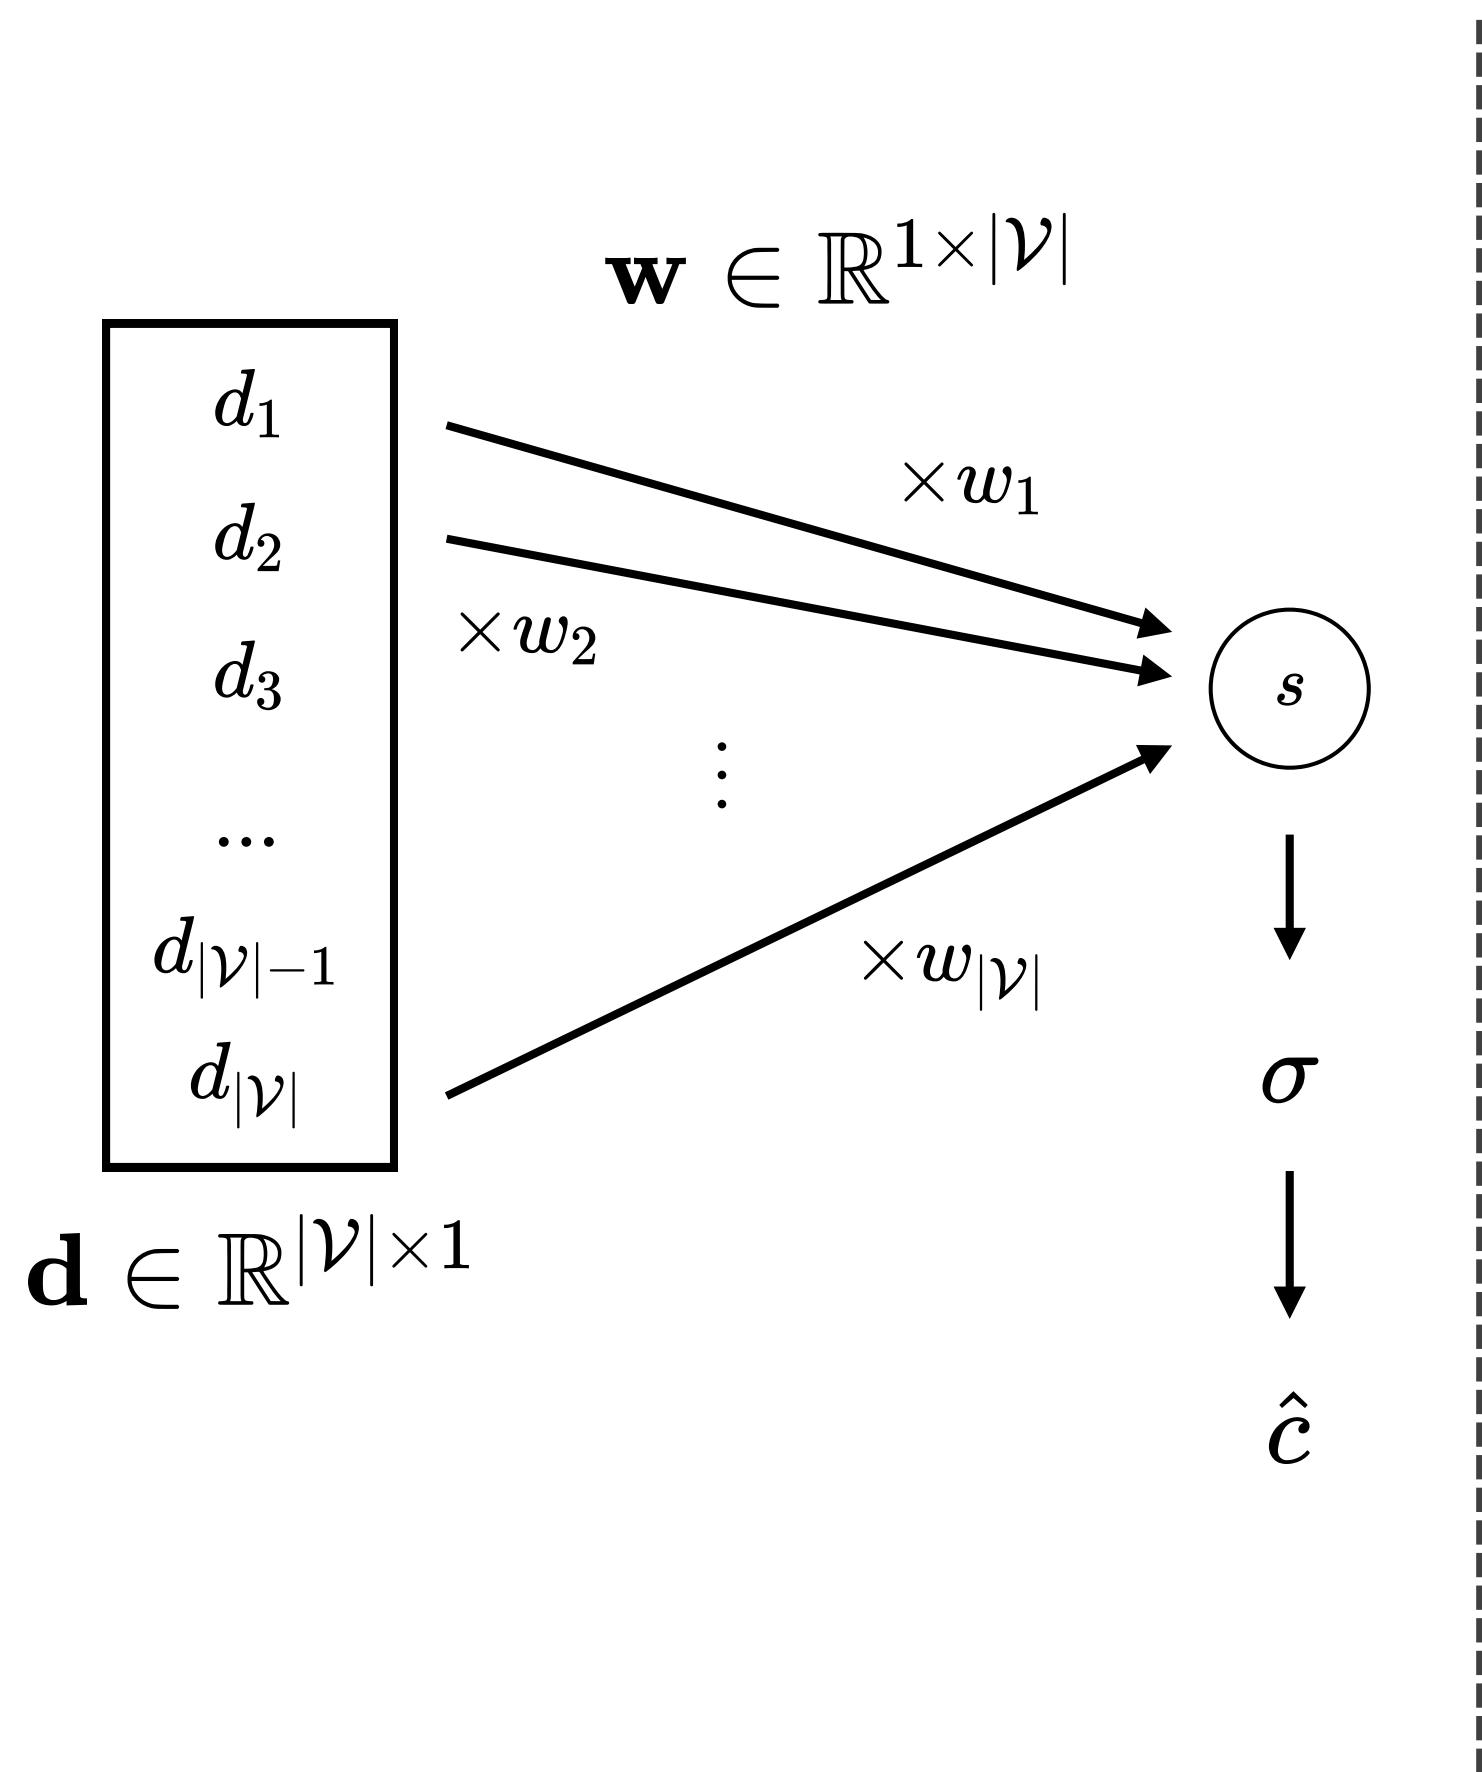
The model is easily extended to a multinomial case through using a matrix  $\mathbf{W}$ , a vector  $\mathbf{b}$  and the *softmax* function (more on that later)



# Classification with logistic regression

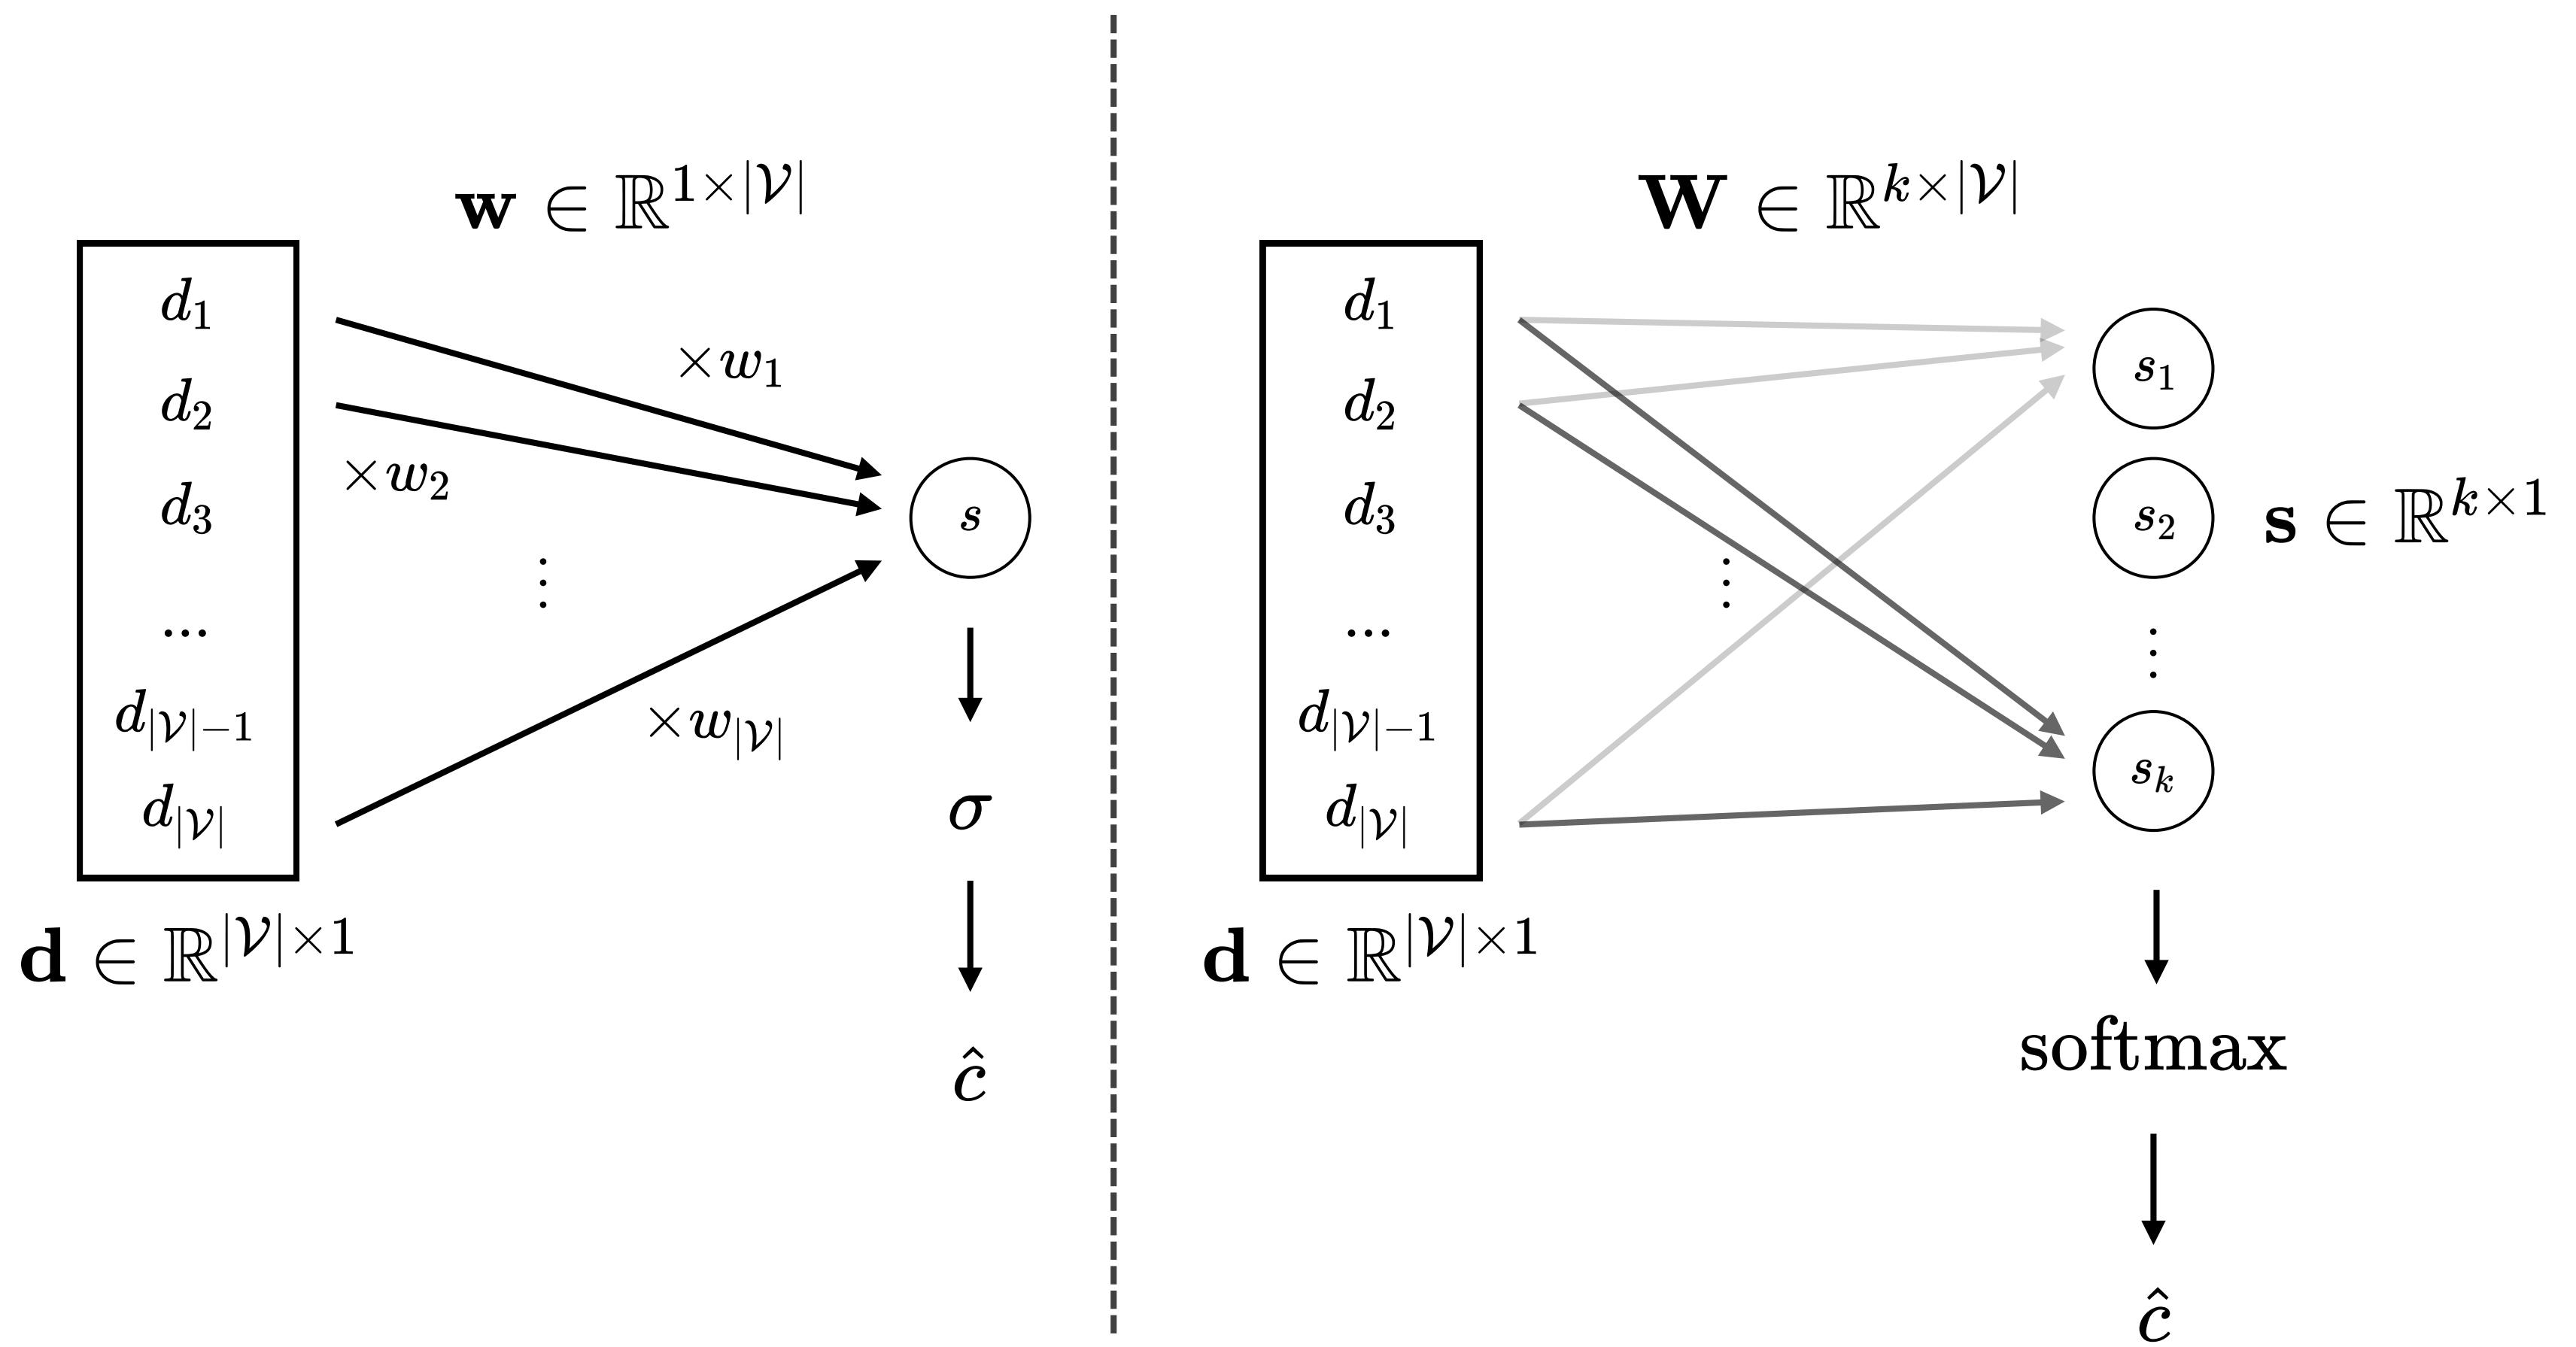
---

The model is easily extended to a multinomial case through using a matrix  $\mathbf{W}$ , a vector  $\mathbf{b}$  and the *softmax* function (more on that later)



# Classification with logistic regression

The model is easily extended to a multinomial case through using a matrix  $\mathbf{W}$ , a vector  $\mathbf{b}$  and the *softmax* function (more on that later)



# Evaluation: Accuracy

---

- As usual: reserve held-out validation set for *hyperparameters tuning* and test set for evaluation
- Simplest measure: **Accuracy**

$$acc(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{|\mathcal{D}|} \sum_{d \in \mathcal{D}} \delta(\hat{y}^d = y^d)$$

- For each label  $c \in \mathcal{C}$ , look at the *type* of
  - Errors: False positive (**FP**) and False negative (**FN**)
  - Correct predictions: True positive (**TP**) and True negative (**TN**)

		Predicted condition	
		Positive (PP)	Negative (PN)
Actual condition	Total population $= P + N$	Positive (P)	Negative (N)
	Positive (P)	True positive (TP)	False negative (FN)
Negative (N)		False positive (FP)	True negative (TN)

# Better metrics: F-measures and Macro

---

- Compute **recall** and **precision**:

$$Recall(\hat{\mathbf{y}}, \mathbf{y}, c) = \frac{TP}{TP + FN} \quad Precision(\hat{\mathbf{y}}, \mathbf{y}, c) = \frac{TP}{TP + FP}$$

- F-measure: combines recall ( $r$ ) and precision ( $p$ ) using the *harmonic mean*

$$\text{F-measure}(\hat{\mathbf{y}}, \mathbf{y}, c) = \frac{2rp}{r + p}$$

- Evaluating **multi-class classification**:

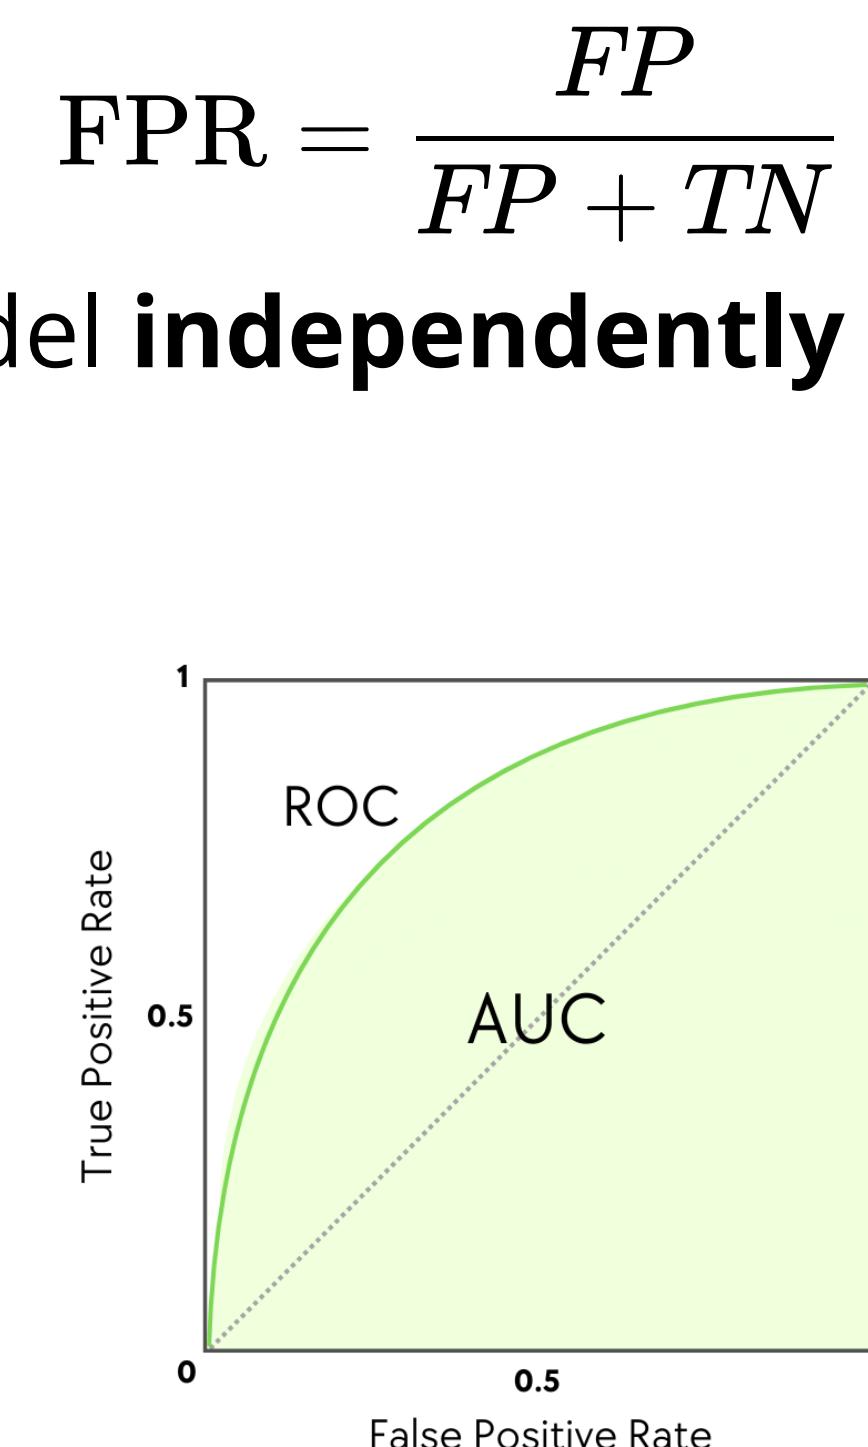
- Balance across instances: Add up **TP**, **FP**, **TN**, **FN** over classes and compute the **Micro** F-measure
- When classes are **imbalanced**, average over classes: **Macro** F-measure:

$$\text{Macro-F}(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{|\mathcal{C}|} \sum_{c \in \mathcal{C}} \text{F-measure}(\hat{\mathbf{y}}, \mathbf{y}, c)$$

# AUC: Area Under the Curve

---

- Area Under the **ROC** (*Receiver Operating Characteristic*) Curve
  - **ROC:**  $\text{TPR}_s = f(\text{FPR}_s)$  for different classification threshold  $s$ 
    - TPR: *True Positive Ratio = Recall*
    - FPR: *False Positive Ratio*
- A metric for measuring the quality of a model **independently from the classification threshold  $s$** 
  - Usually,  $s = \frac{1}{2}$
  - Can be adapted to unbalanced tasks  
→ Anomaly detection
- The higher the AUC, the better



# Back to difficulty 2: Ambiguity

---

To go further, NLP systems usually need to *uncover the **structure*** of text, which is made difficult by:

- **Lexical ambiguity:** *homography, polysemy*
- **Syntactic ambiguity**
- Ambiguity in **semantic scope**

# Back to difficulty 2: Ambiguity

---

To go further, NLP systems usually need to *uncover the **structure*** of text, which is made difficult by:

- **Lexical ambiguity:** *homography, polysemy*
  - Task: Part-of-speech tagging, Word sense disambiguation
- **Syntactic ambiguity**
  - Task: Dependency parsing
- Ambiguity in **semantic scope**
  - Task: Semantic parsing

# Back to difficulty 2: Ambiguity

---

To go further, NLP systems usually need to *uncover the **structure*** of text, which is made difficult by:

- **Lexical ambiguity:** *homography, polysemy*
  - Task: Part-of-speech tagging, Word sense disambiguation
- **Syntactic ambiguity**
  - Task: Dependency parsing
- Ambiguity in **semantic scope**
  - Task: Semantic parsing

Again, things are complicated by lacking *implicit knowledge*:

- Background, commonsense knowledge
- Contextual knowledge

# Features: *Implicit* representations

Beyond being counted, words can be represented by **explicit** features:

- List of **attributes** describing the object
- Natural language definition (*dictionary*)
- Other lexical resources, including senses and associated properties, morphological features: *WordNet*

WordNet Search - 3.1  
- [WordNet home page](#) - [Glossary](#) - [Help](#)

Word to search for:

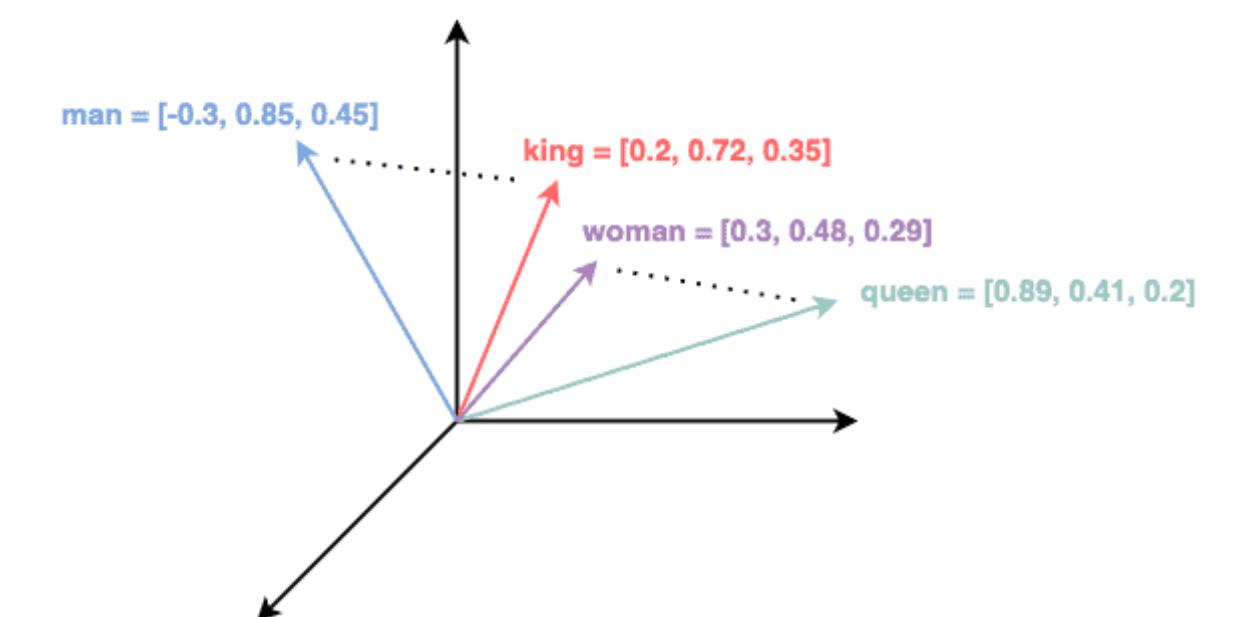
Display Options:    
Key: "S:" = Show Synset (semantic) relations, "W:" = Show Word (lexical) relations  
Display options for sense: (gloss) "an example sentence"

Noun

- [S: \(n\) wordnet](#) (any of the machine-readable lexical databases modeled after the Princeton WordNet)
- [S: \(n\) WordNet, Princeton WordNet](#) (a machine-readable lexical database organized by meanings; developed at Princeton University)

→ Move to *implicit* features: with **distributed representations**

- **Vector spaces** for words: encode **contextual** information
  - *Distributional hypothesis*: two words are similar if they have similar contexts
  - Create sparse then dense representations → Embeddings



# Model: using context

---

Work on modelisation of the **context** - which is, most of the time, the *surrounding sequence*:

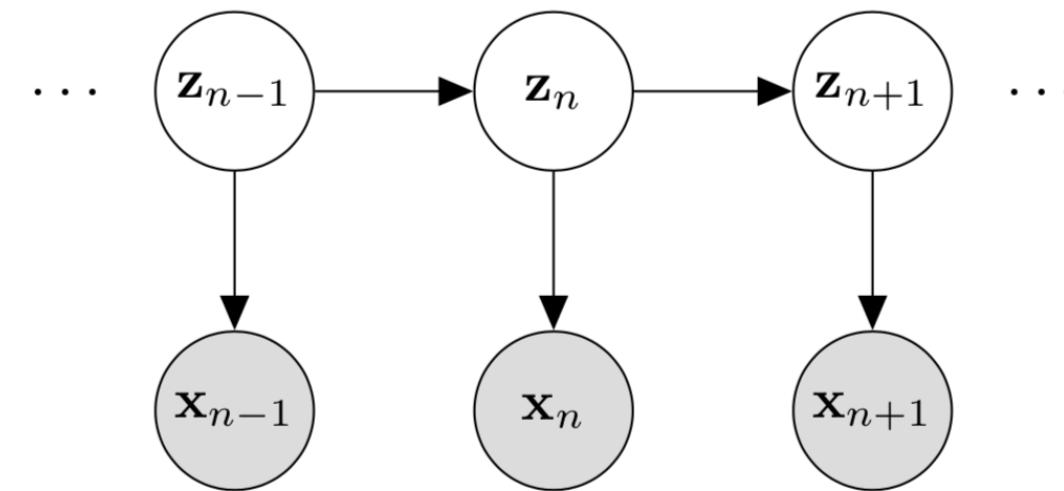
- Use appropriate **sequence models** which will get the necessary information from the immediate context

# Model: using context

---

Work on modelisation of the **context** - which is, most of the time, the *surrounding sequence*:

- Use appropriate **sequence models** which will get the necessary information from the immediate context
- Model **dependency within the sequence**: *Markov models*
- Generative modeling:
  - *Hidden Markov Models*



# Model: using context

---

Work on modelisation of the **context** - which is, most of the time, the *surrounding sequence*:

- Use appropriate **sequence models** which will get the necessary information from the immediate context
- Model **dependency within the sequence**: *Markov models*
- Generative modeling:
  - *Hidden Markov Models*
- In this class: *Deep learning sequential models*
  - Architectures and objectives designed to take advantage of **large-scale unlabelled datasets**
  - Interaction with **traditional tasks, structures** ?
  - Integration of **exterior knowledge** ?

